

# Software Defined Network

7COM1030 – Multicast and Multimedia Networking

Dr. Xianhui Che (Cherry)

[x.che@herts.ac.uk](mailto:x.che@herts.ac.uk)

School of Engineering and Computer Science  
University of Hertfordshire, UK

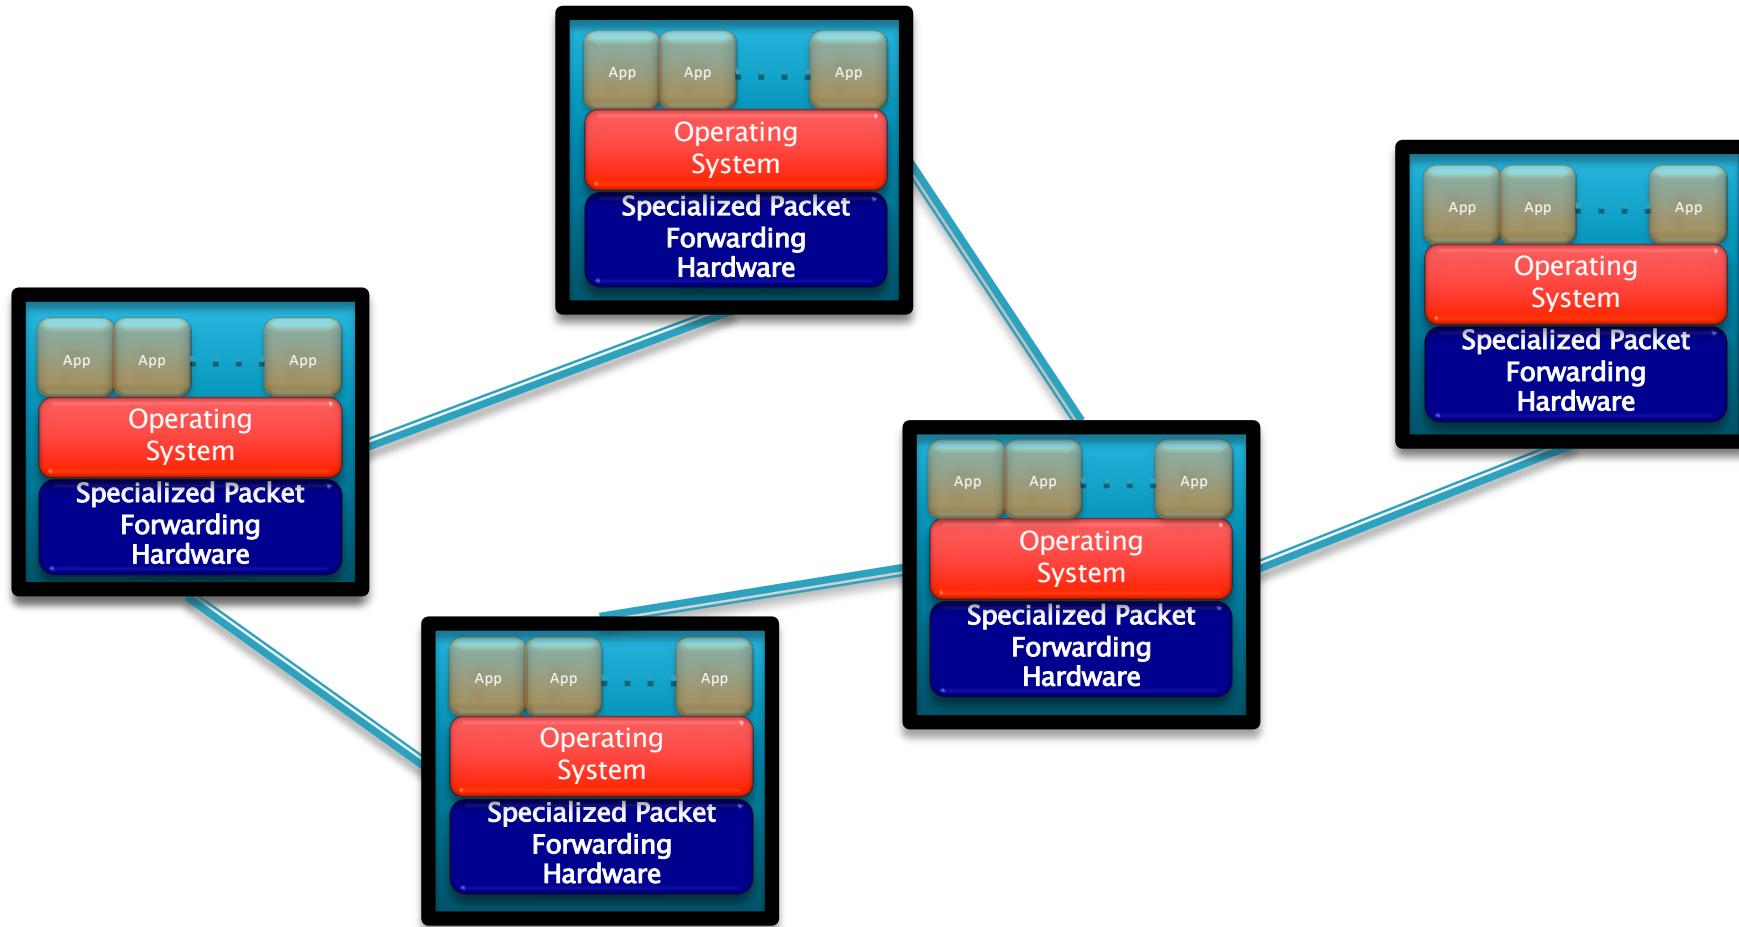


# Topics

- ▶ Needs for SDN
- ▶ SDN Principles
- ▶ OpenFlow
- ▶ Impact on Future Networking

# What is wrong with the legacy Internet?

Closed to innovations in the infrastructure



# The Needs for SDN (1)

## ► 1. New traffic patterns

- **Traffic is no longer simply end-to-end.** In contrast to client-server applications where the bulk of the communication occurs between one client and one server, today's applications access different databases and servers, creating a flurry of machine-to-machine traffic before returning data to the end user device.
- **Ubiquitous access.** Users are changing network traffic patterns as they push for access to corporate content and applications from any type of device (including their own), connecting from anywhere, at any time.
- **Additional traffic for cloud computing.** Many enterprise data centre managers are contemplating a utility computing model, which might include a private cloud, public cloud, or some mix of both, resulting in additional traffic across the wide area network.

# The Needs for SDN (2)

## ► 2. The network is hard to evolve

- Ongoing innovation in systems software, e.g. new languages, operating systems, etc.
- However -
  - Networks are stuck in the past
  - Routing algorithms change very slowly
  - Network management extremely primitive

# The Needs for SDN (3)

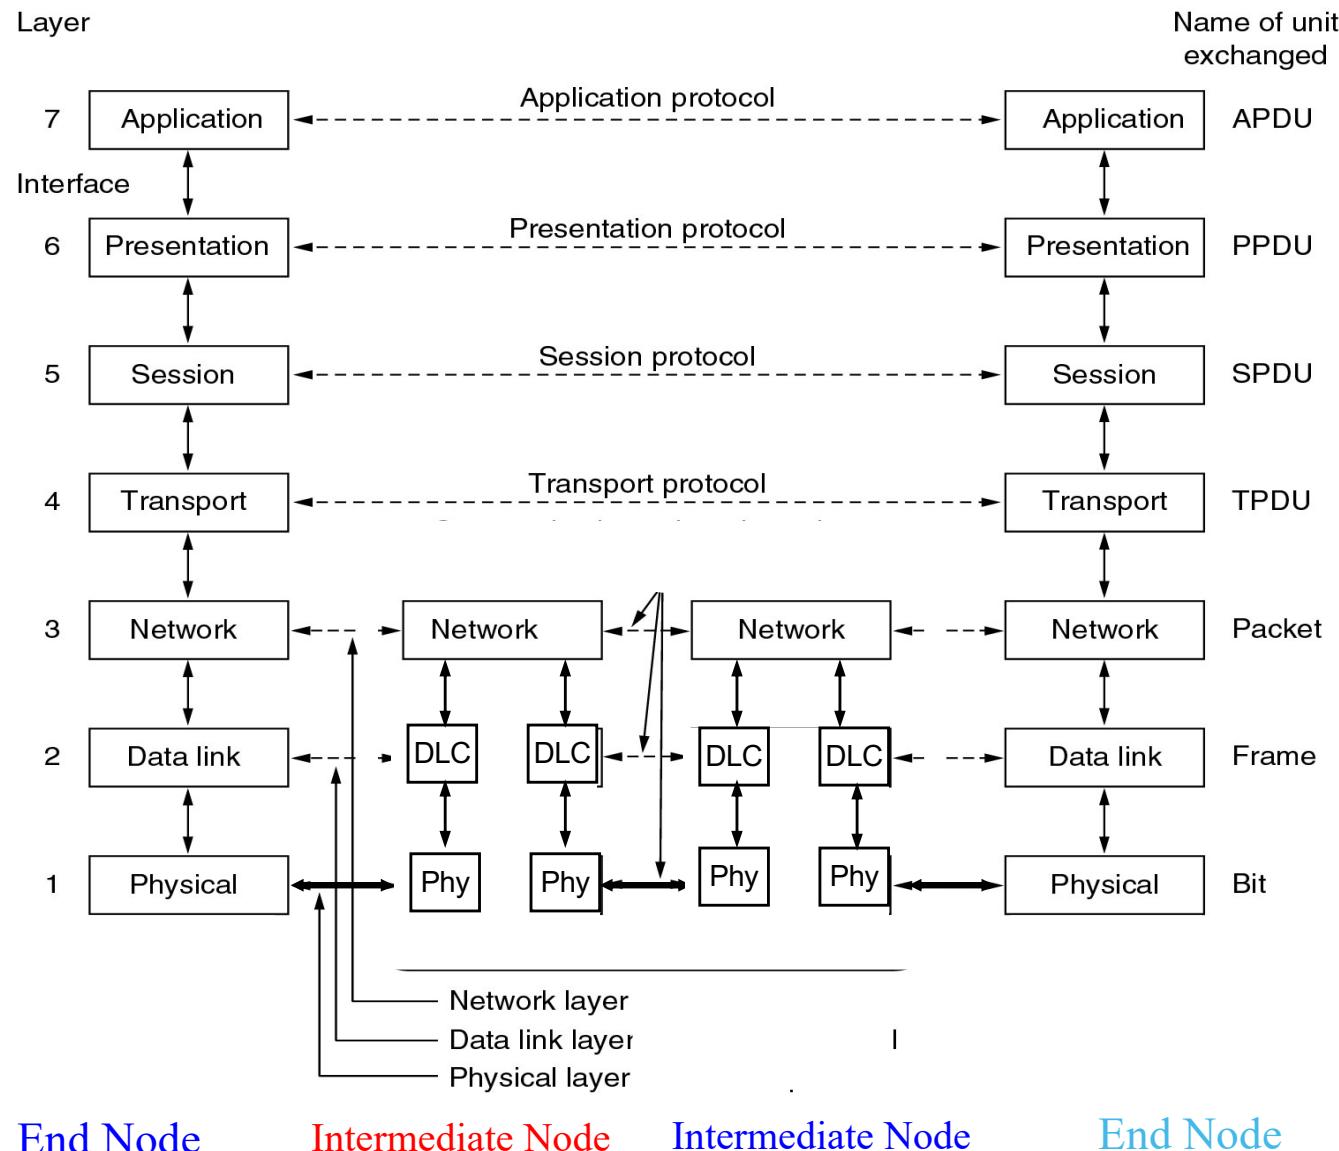
- ▶ **3. No formal principles on network design**
  - OS courses teach fundamental principles
  - Networking courses teach a big bag of protocols
  - No formal principles, just general design guidelines

# The Needs for SDN (4)

## ► 4. Emerging architectures

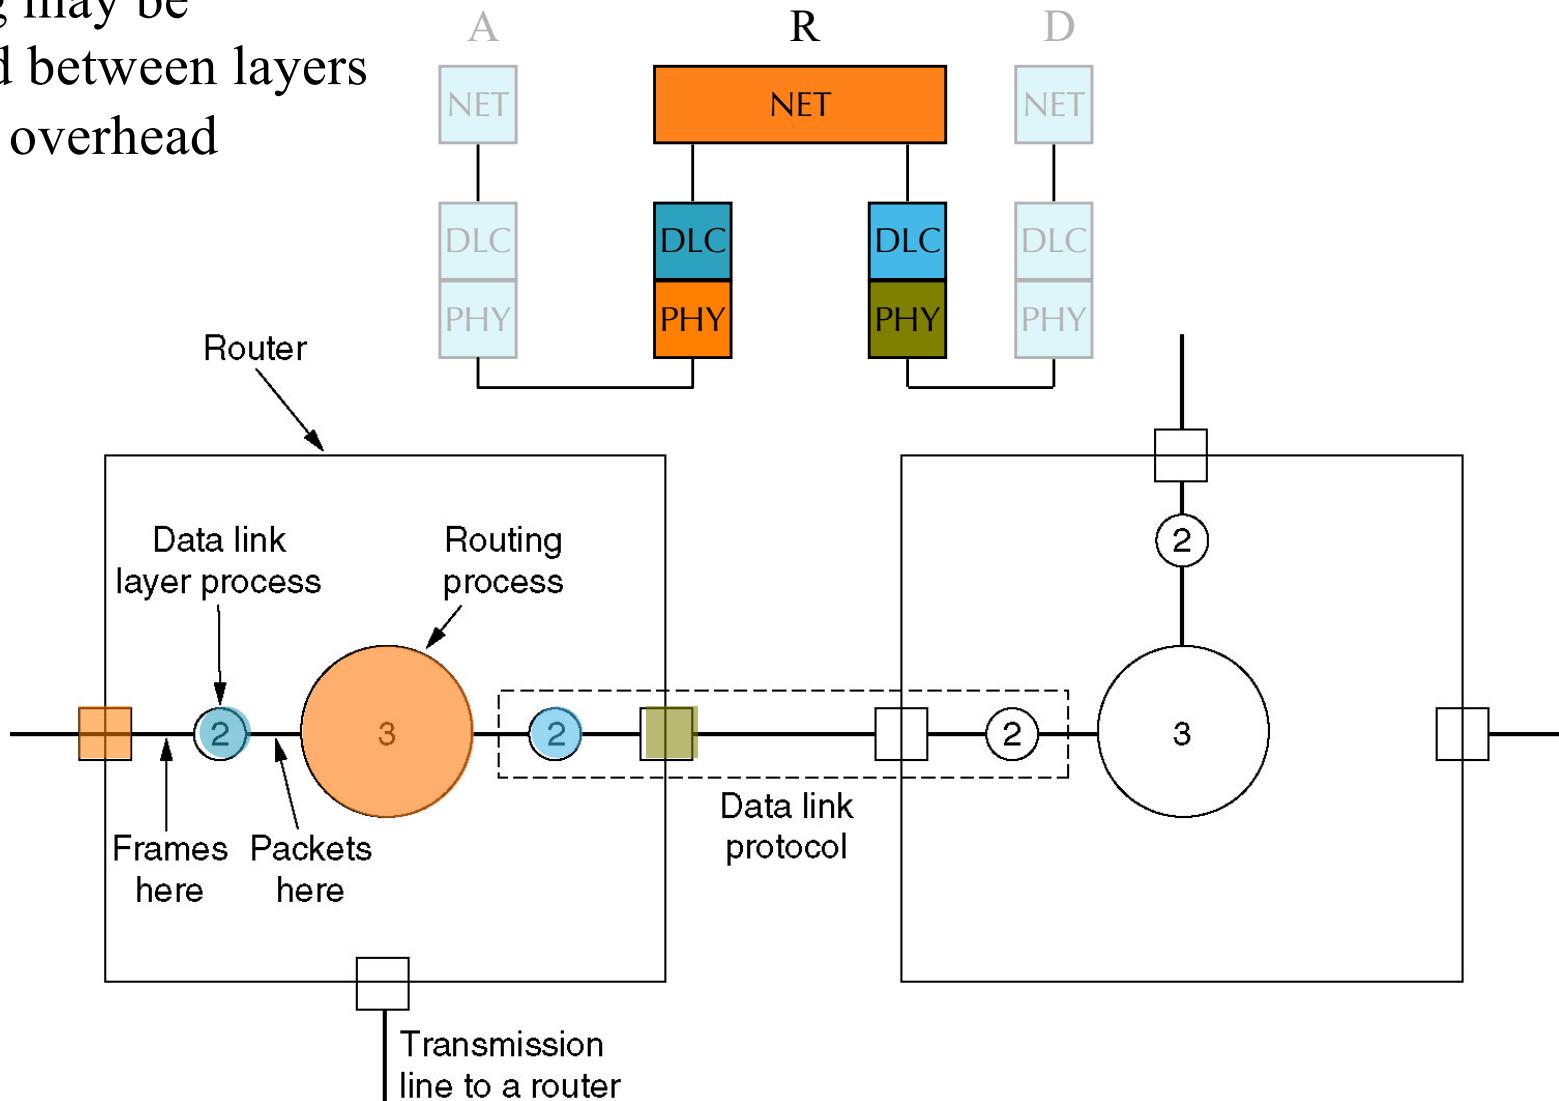
- **Mobile computing model.** Users are increasingly employing mobile personal devices such as smartphones, tablets, and notebooks to access the corporate network. IT is under pressure to accommodate these personal devices in a fine-grained manner while protecting corporate data and intellectual property and meeting compliance mandates.
- **Cloud computing services.** Enterprises have enthusiastically embraced both public and private cloud services, resulting in unprecedented growth of these services. Providing self-service provisioning, whether in a private or public cloud, requires elastic scaling of computing, storage, and network resources.

# Peer Processes

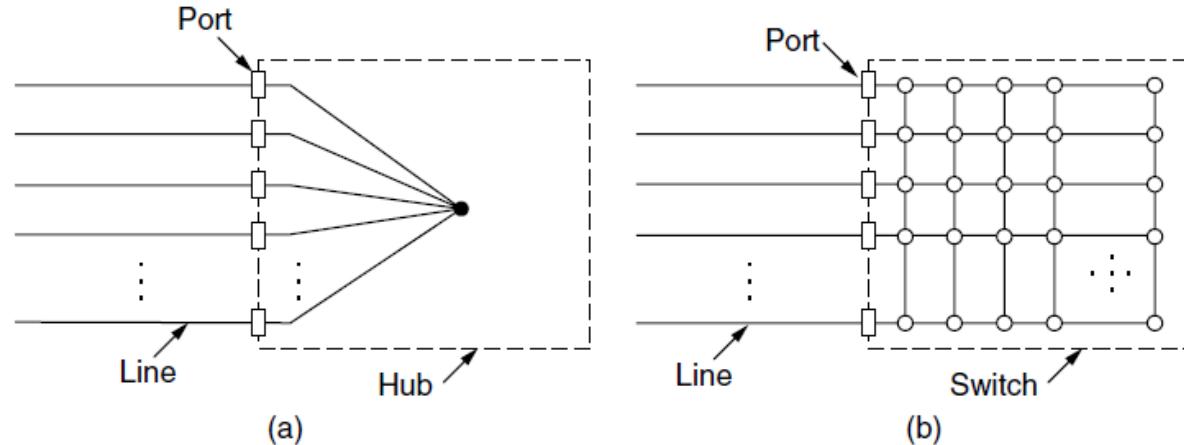


# Inter-Protocol Operation

- ▶ Buffering may be employed between layers  
→ added overhead latency



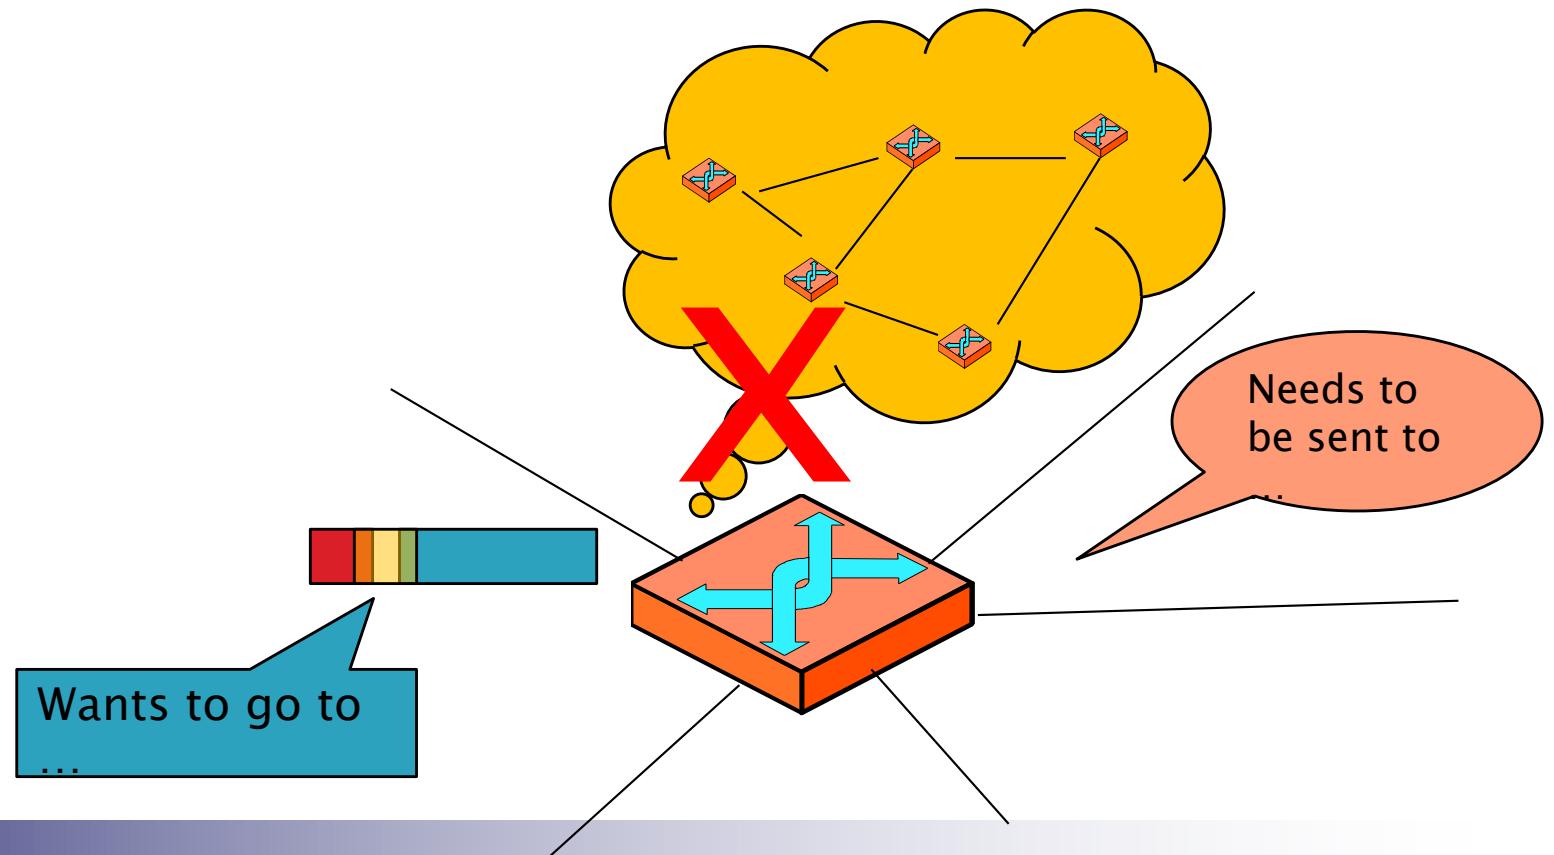
# Switched Ethernet



- ▶ “Hub” contracted long wire into a box
- ▶ “Switch” replaces short wire in box with a switchable network → **Switching is faster than routing!**
  - No collisions – no MAC!
- ▶ Supports simultaneous transmissions
  - But needs buffering

# Forwarding Decision

- ▶ Made by each forwarding (intermediate) node (“router”) for the next hop only
  - (Why just the next hop, just ”one-hop” forwarding?)



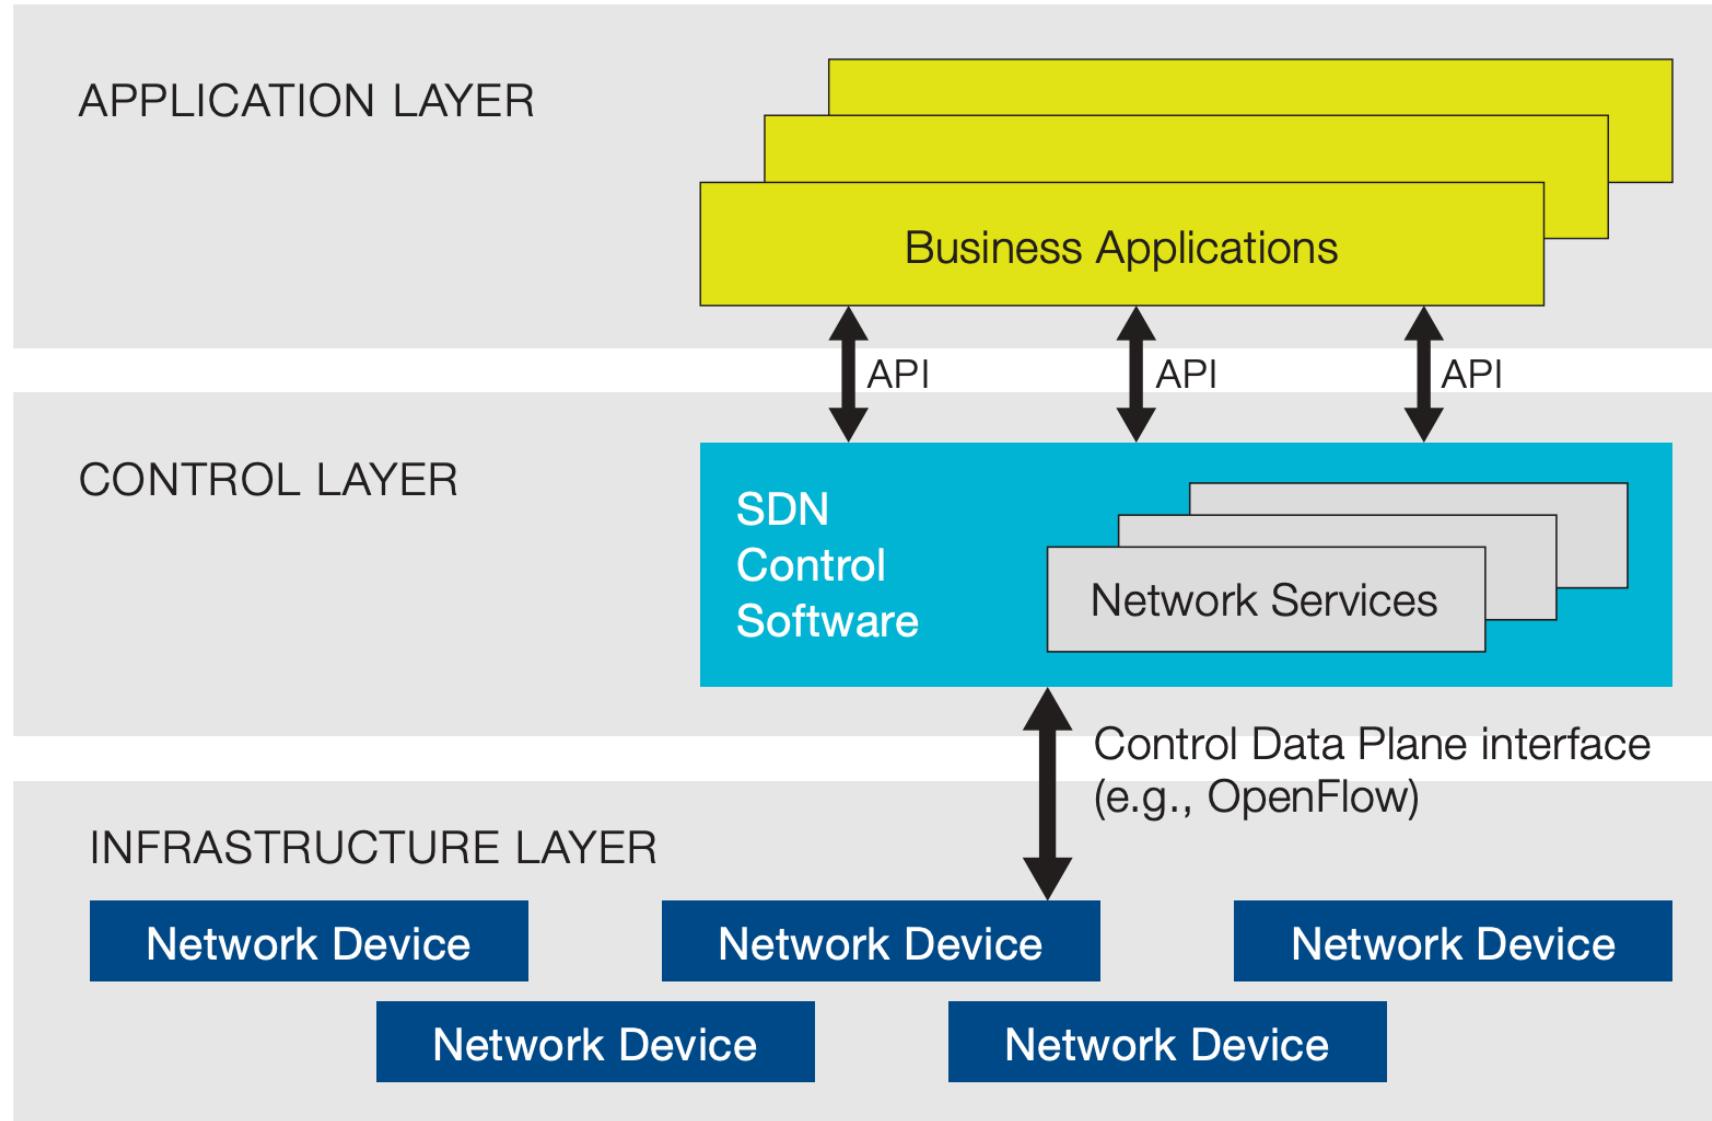
# Implications

- ▶ Networking is an exercise in coordination
  - Local actions must coordinate to produce global effect
  - Agreement on global effect desired is essential → standardized protocols
- ▶ Local actions are constrained to produce same global effect
- ▶ Changing global effect agreement requires:
  - Agreement, if multi-owner
  - Changing local action globally across system

# Topics

- ▶ Needs for SDN
- ▶ SDN Principles
- ▶ OpenFlow
- ▶ Impact on Future Networking

# SDN Architecture



# SDN: Two Control Plane Abstractions

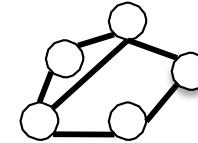
- Abstraction: **global network view**
  - Provides information about current network
  - **Implementation:** “Network Operating System”
    - Runs on servers in network (replicated for reliability)
- Abstraction: **forwarding model**
  - Provides standard way of defining forwarding state
  - This is OpenFlow
    - Specification of <match,action> flow entries

# Networks of Software Defined Routers

routing, access control, etc.

Control Program

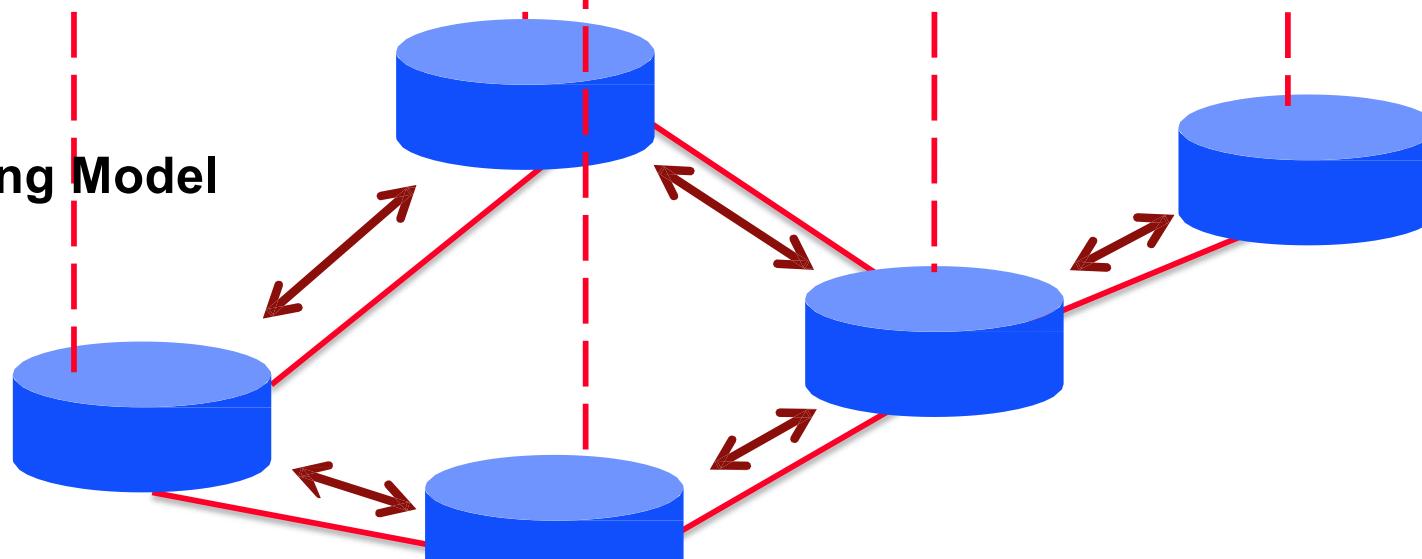
Global Network View



Distributed Neighbors

Network OS (e.g. NOX)

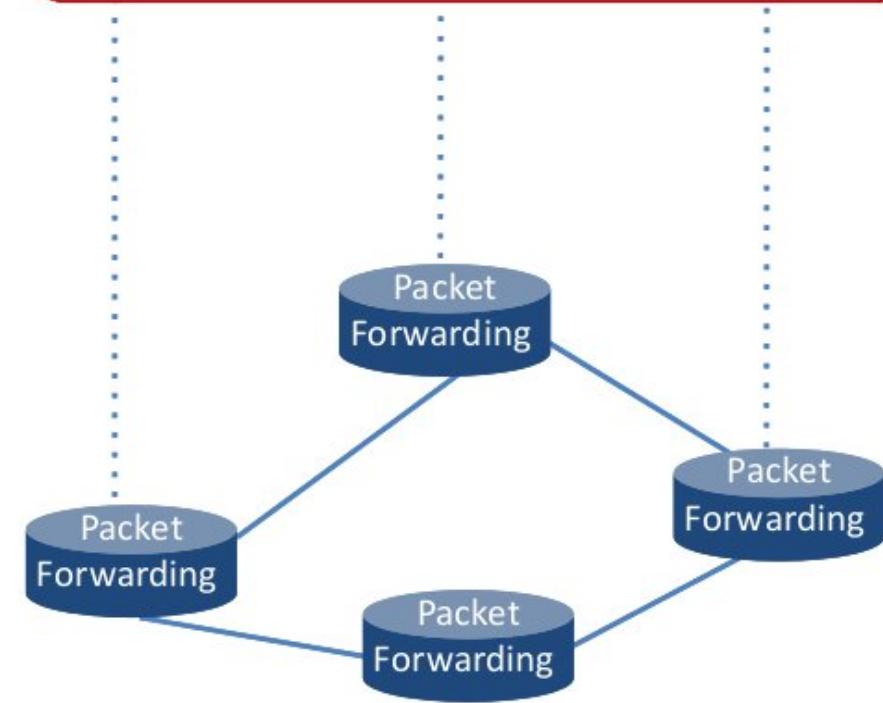
Forwarding Model



# Example: OSPF and Dijkstra

- OSPF
  - RFC 2328: 245 pages
- Distributed System
  - Builds consistent, up-to-date map of the network: 101 pages
- Dijkstra's Algorithm
  - Operates on map: 4 pages

# Example: OSPF and Dijkstra

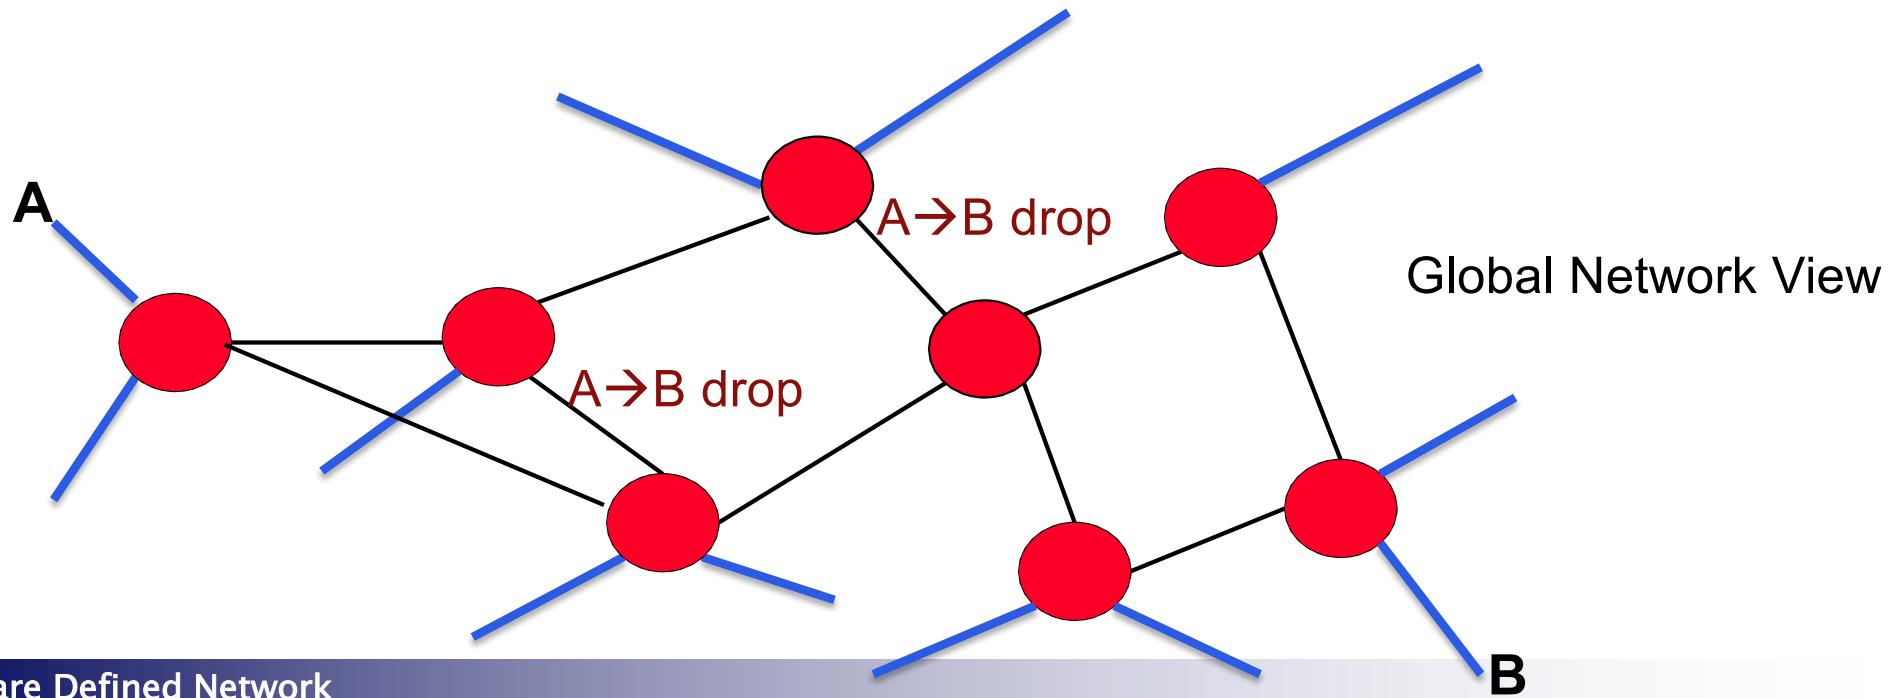


# Specification Abstraction

- Control program must **express** desired behavior
  - Whether it be isolation, access control, or QoS
- It should not be responsible for **implementing** that behavior on physical network infrastructure
  - Requires configuring the forwarding tables in each switch
- Proposed abstraction: **Virtual Topology** of network
  - Virtual Topology models only enough detail to specify goals
  - Will depend on task semantics

# Simple Example: Access Control

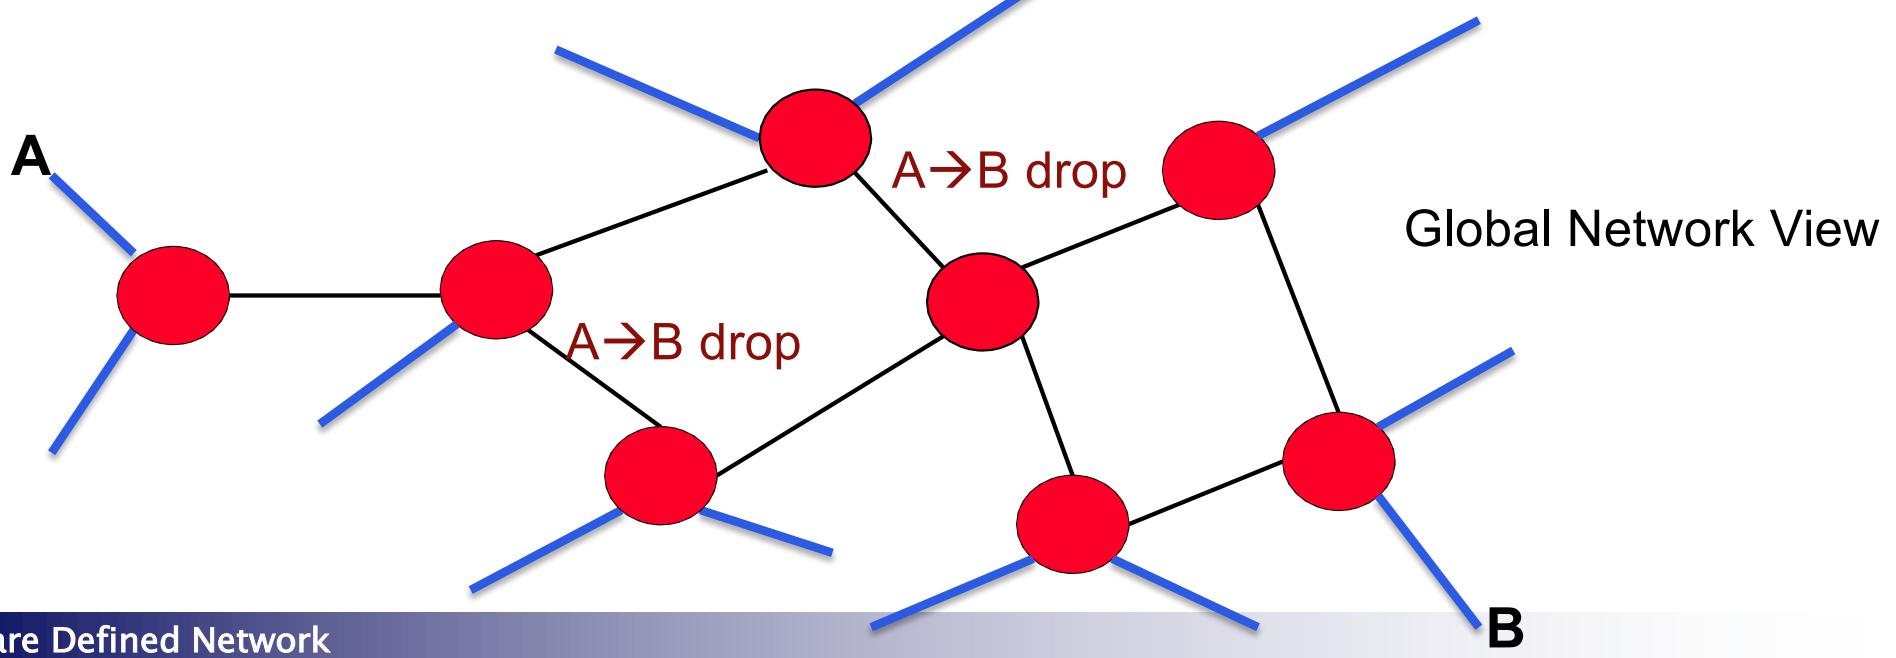
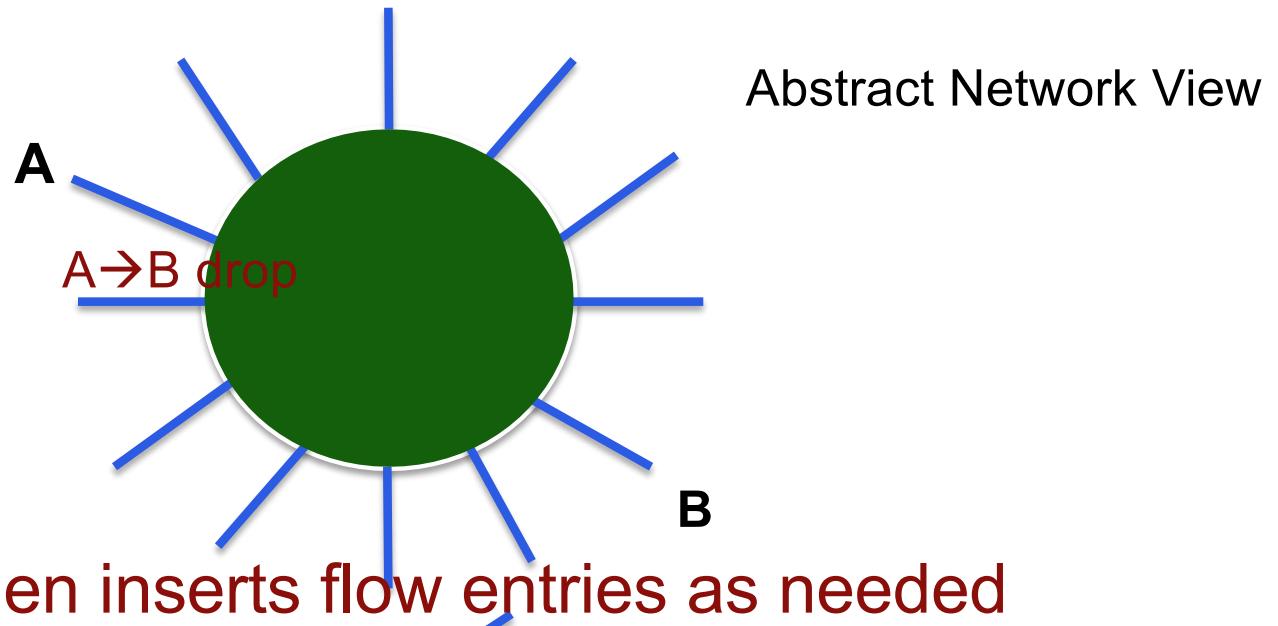
- Operator's goal: prevent A's packets from reaching B
- Control program does so with access control entries:
  - Control program must respond to topology/routing changes
  - Makes it hard to write correct control program



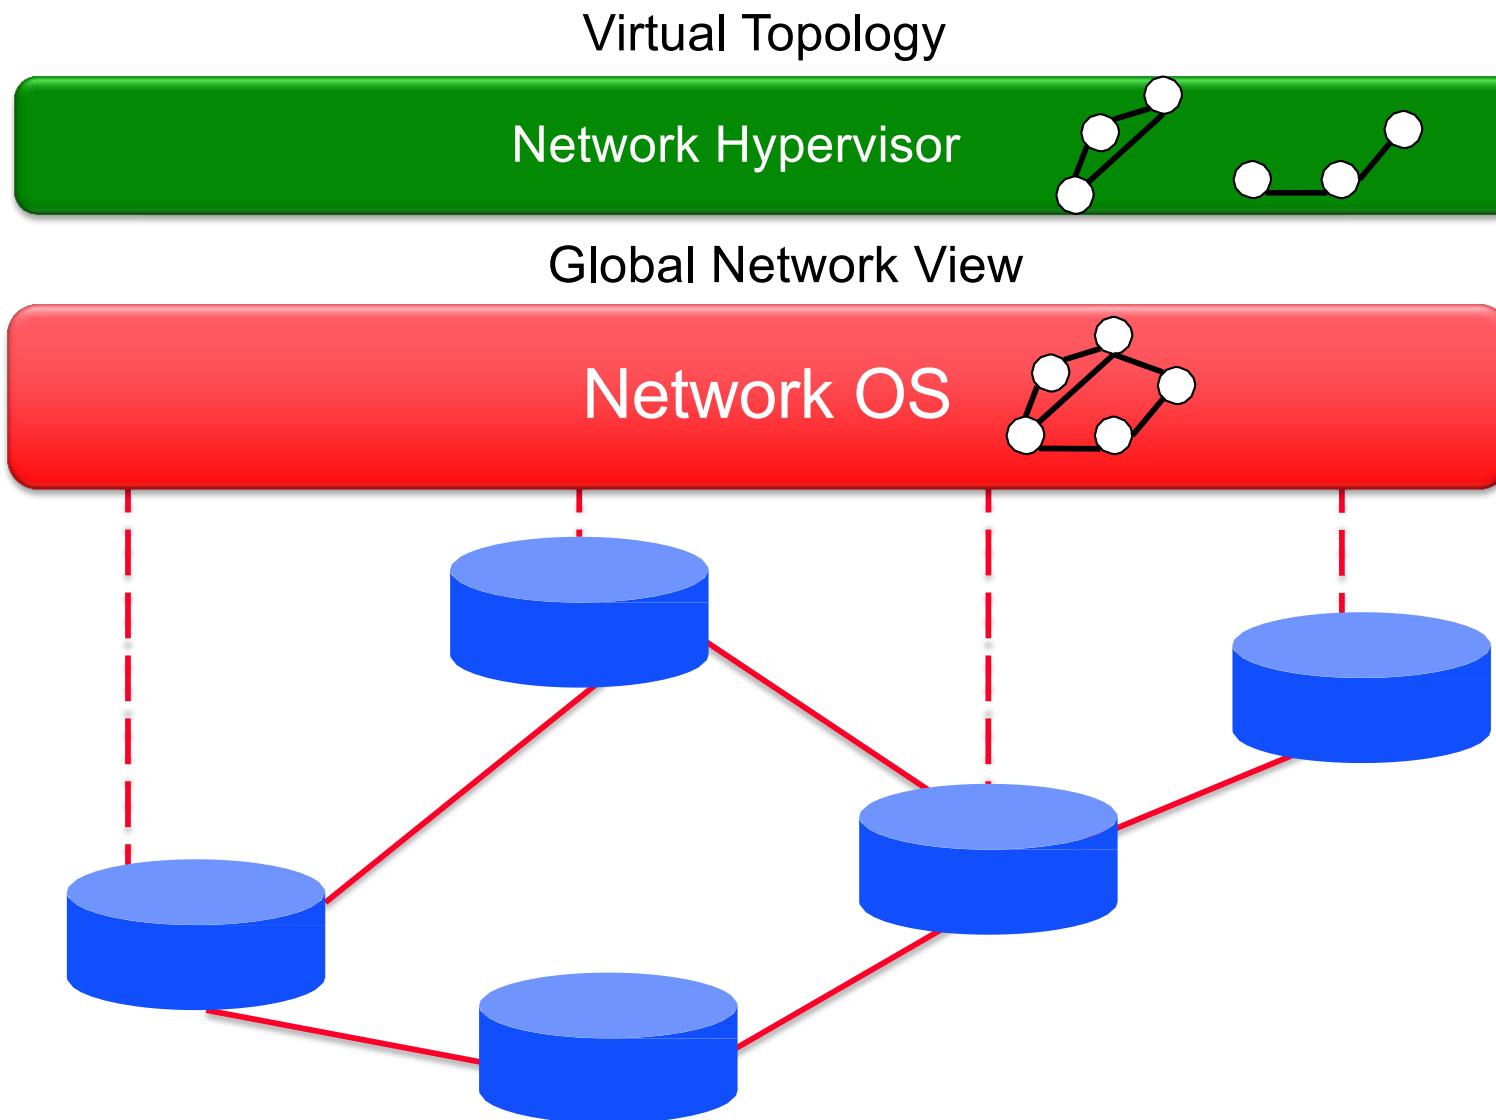
# Network Virtualization

- Introduce new abstraction and new SDN layer
- Abstraction: **Virtual Topology**
  - Allows operator to express requirements and policies
  - Via a set of logical switches and their configurations
- Layer: **Network Hypervisor**
  - Translates those requirements into switch configurations
  - “Compiler” for virtual topologies

# Virtualization Simplifies Control Program



# Software Defined Network Model



# SDN Elements

- **Control program:** express goals on Virtual Topology
  - Operator Requirements
  - Configuration = Function(view)
  - Not a distributed protocol, now just a graph algorithm
- **Network Hypervisor:** Virtual Topology  $\leftrightarrow$  Global Network View
- **Network OS:** Global Network View  $\leftrightarrow$  physical switches
  - Gathers information for global network view
  - Conveys configurations from control program to switches
- **Router/switches:** merely follow orders from NOS

## Clean separation of control and data planes

- Not packaged together in proprietary boxes; enables the use of commodity hardware, 3rd party software
- Easier to write, maintain, and verify

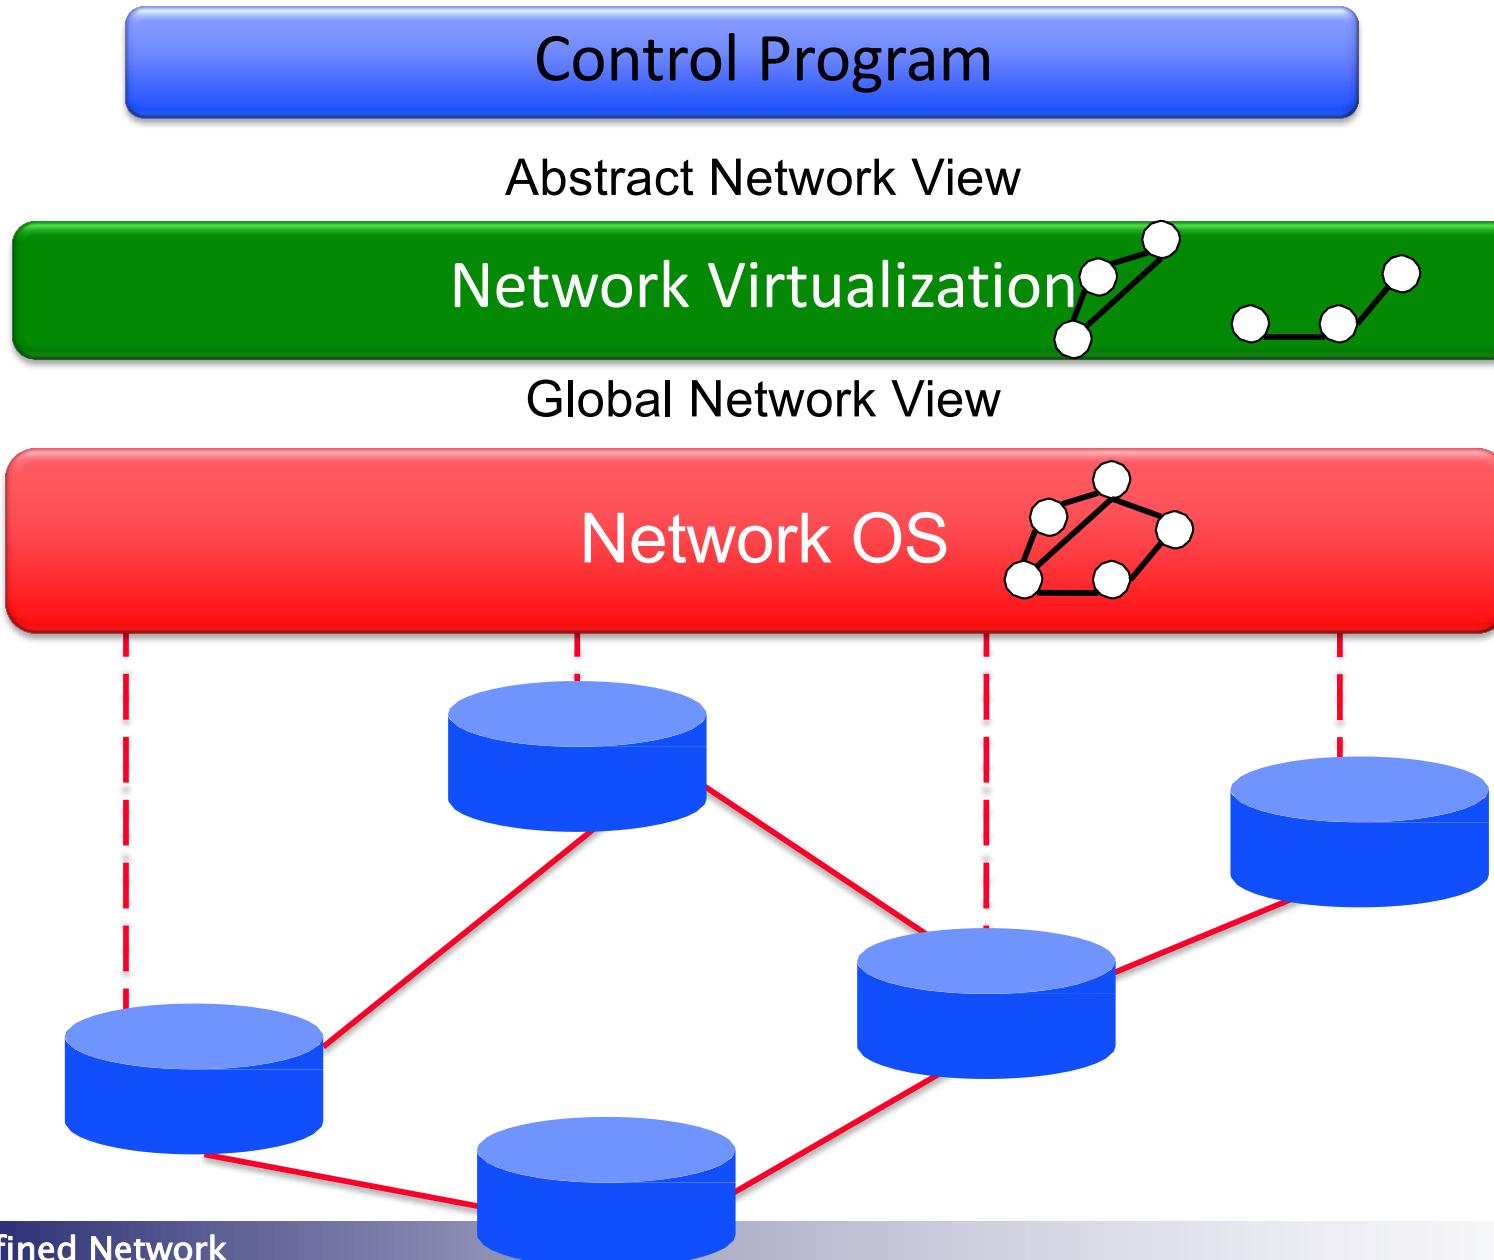
# Abstractions Don't Eliminate Complexity

- Every component of system is tractable
  - NOS, Virtualization are still complicated pieces of code
- SDN main achievements:
  - Simplifies interface for control program (user-specific)
  - Pushes complexity into reusable code (SDN platform)
- Just like compilers....

# Virtualization is Killer App for SDN

- Consider a multi-tenant data center
  - Want to allow each tenant to specify virtual topology
  - This defines their individual policies and requirements
- Data center's network hypervisor compiles these virtual topologies into set of switch configurations
  - Takes 1000s of individual tenant virtual topologies
  - Computes configurations to implement all simultaneously
- ***This is what people are paying money for....***
  - ***Enabled by SDN's ability to virtualize the network***

# SDN: Layers for the Control Plane



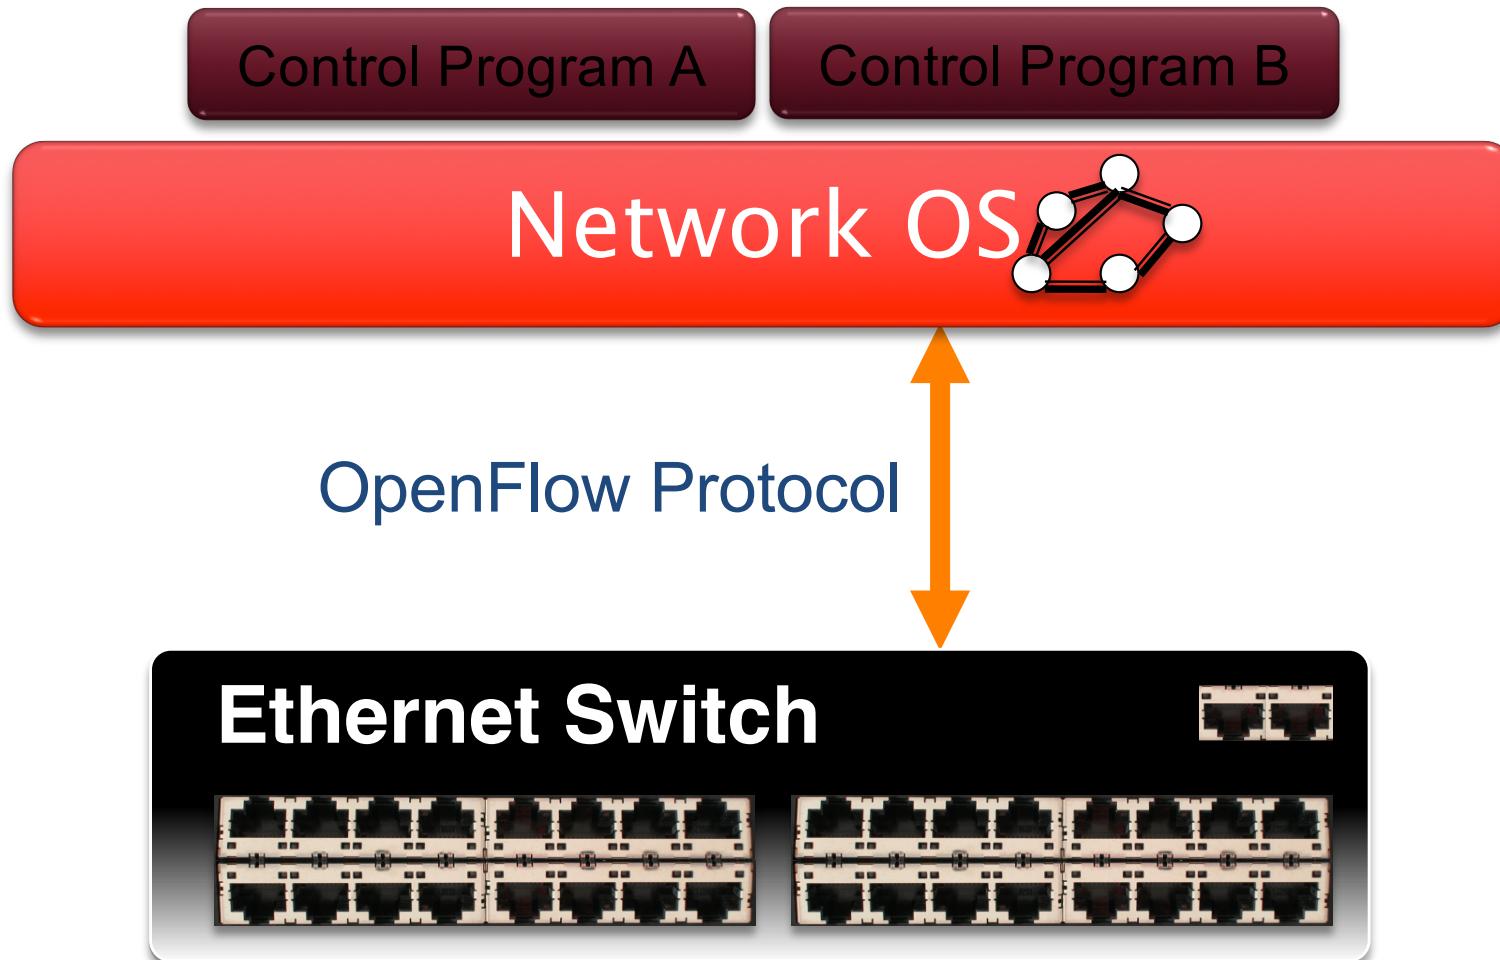
## Four Crucial Points about SDN

- SDN is merely set of abstractions for control plane
  - Not a specific set of mechanisms
  - OpenFlow is least interesting aspect of SDN, technically
- SDN involves computing a function....
  - NOS handles distribution of state
- ...on an abstract network
  - Can ignore actual physical infrastructure
- Network virtualization is the “killer app”
  - Already virtualized compute, storage; network is next

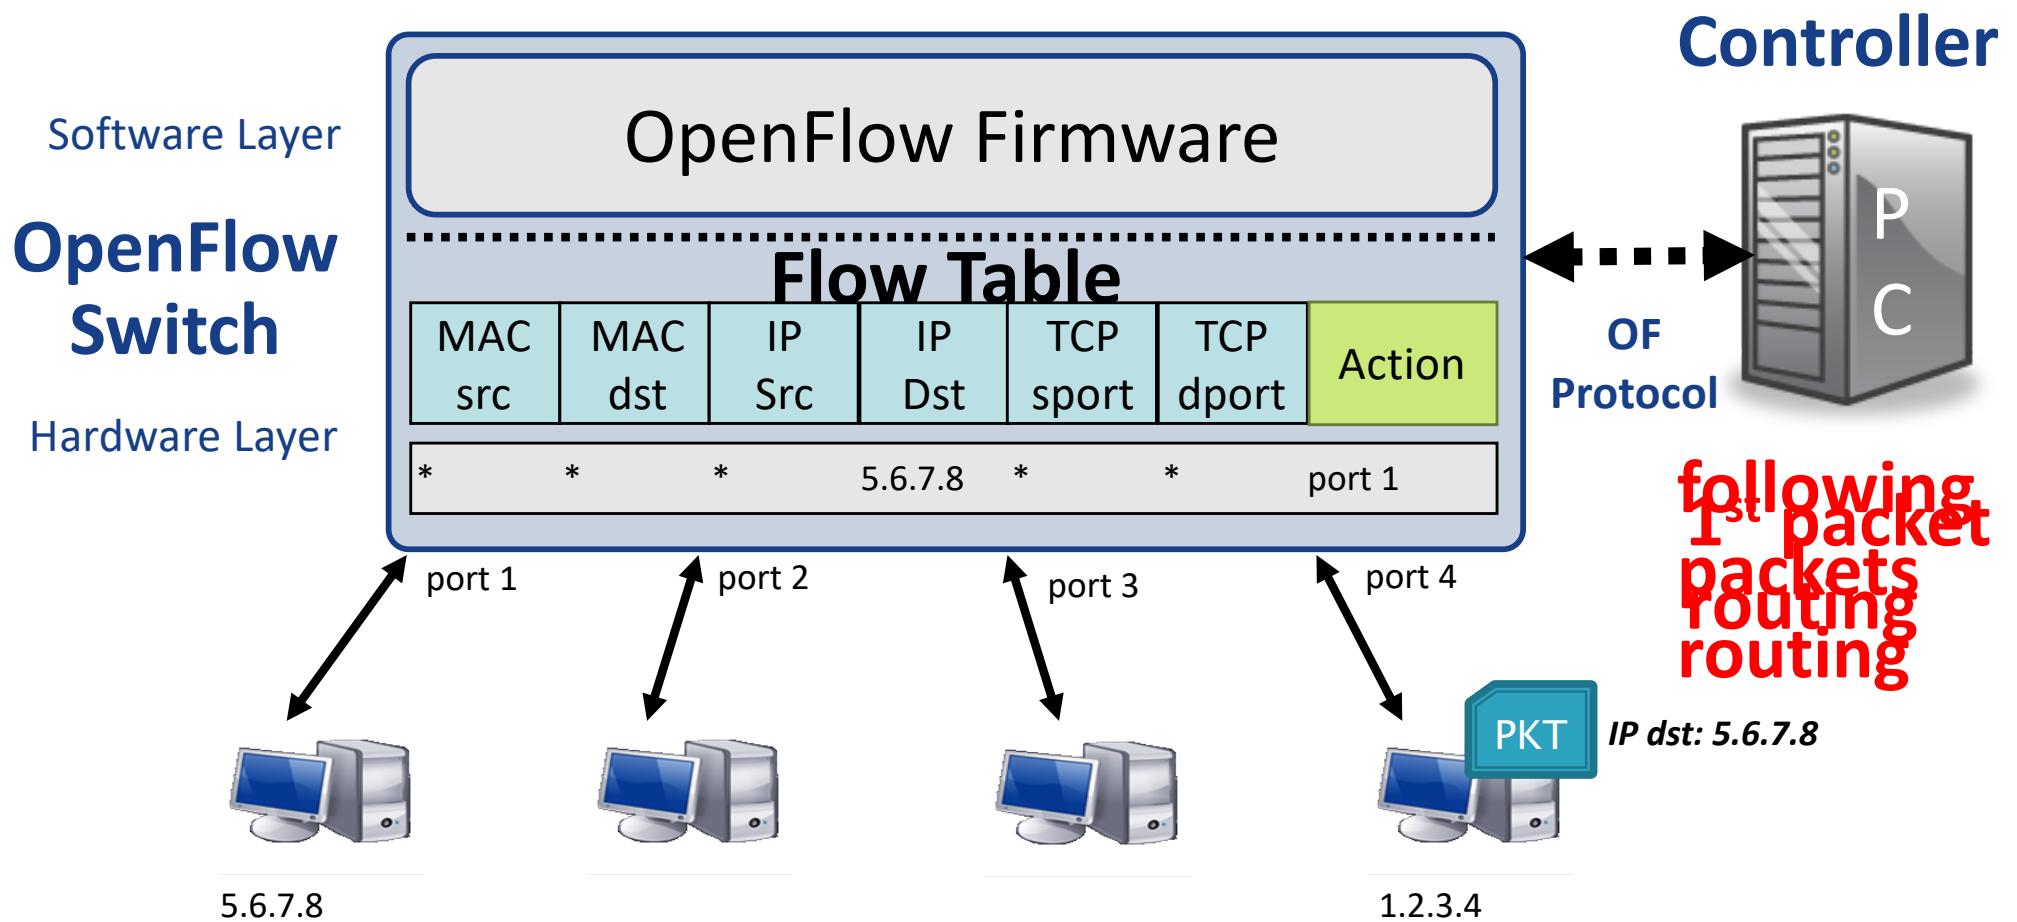
# Topics

- ▶ Needs for SDN
- ▶ SDN Principles
- ▶ OpenFlow
- ▶ Impact on Future Networking

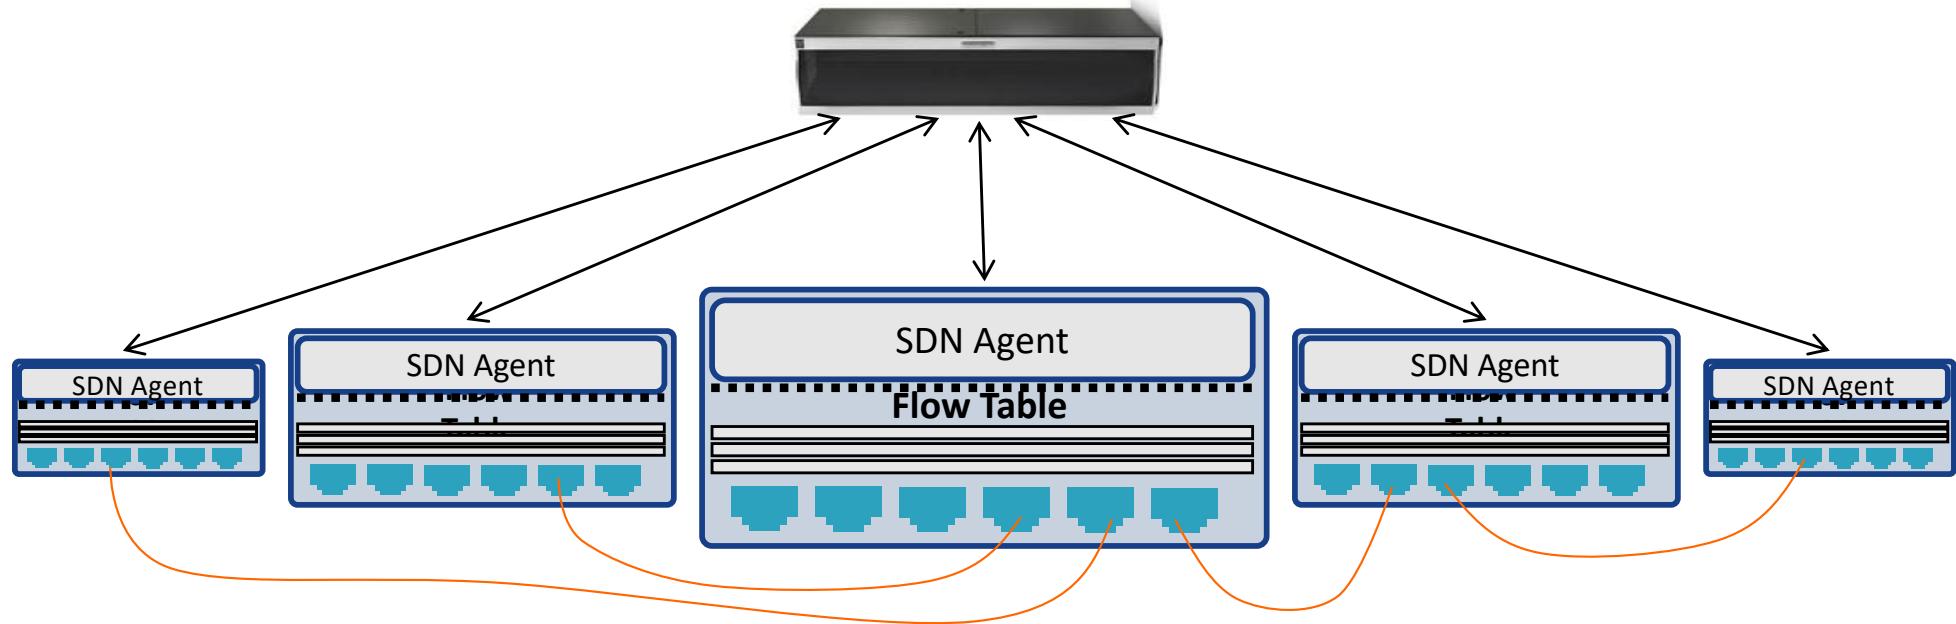
# OpenFlow Basics



# OpenFlow Example

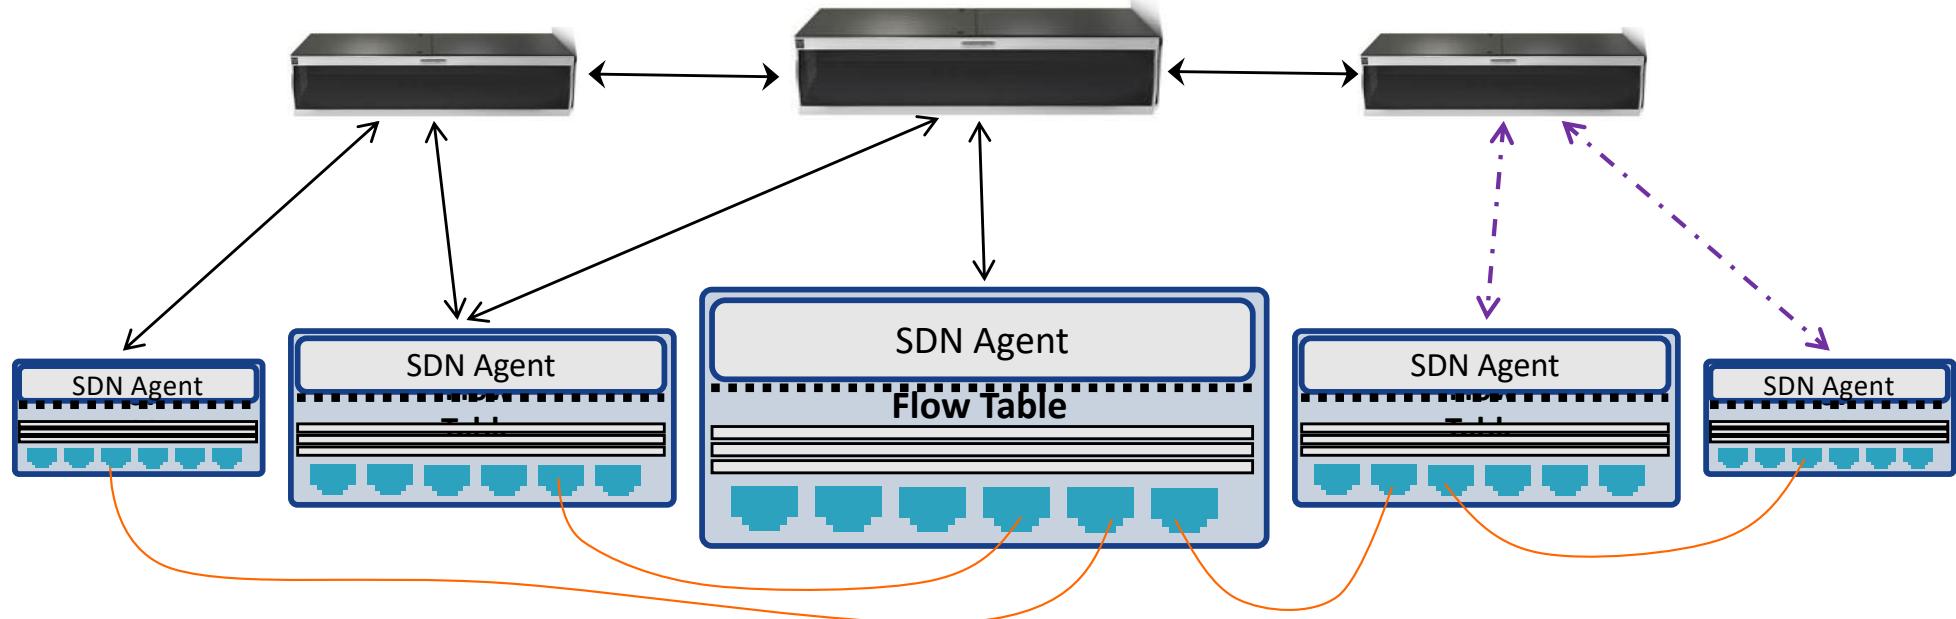


# Controller – Policy Box



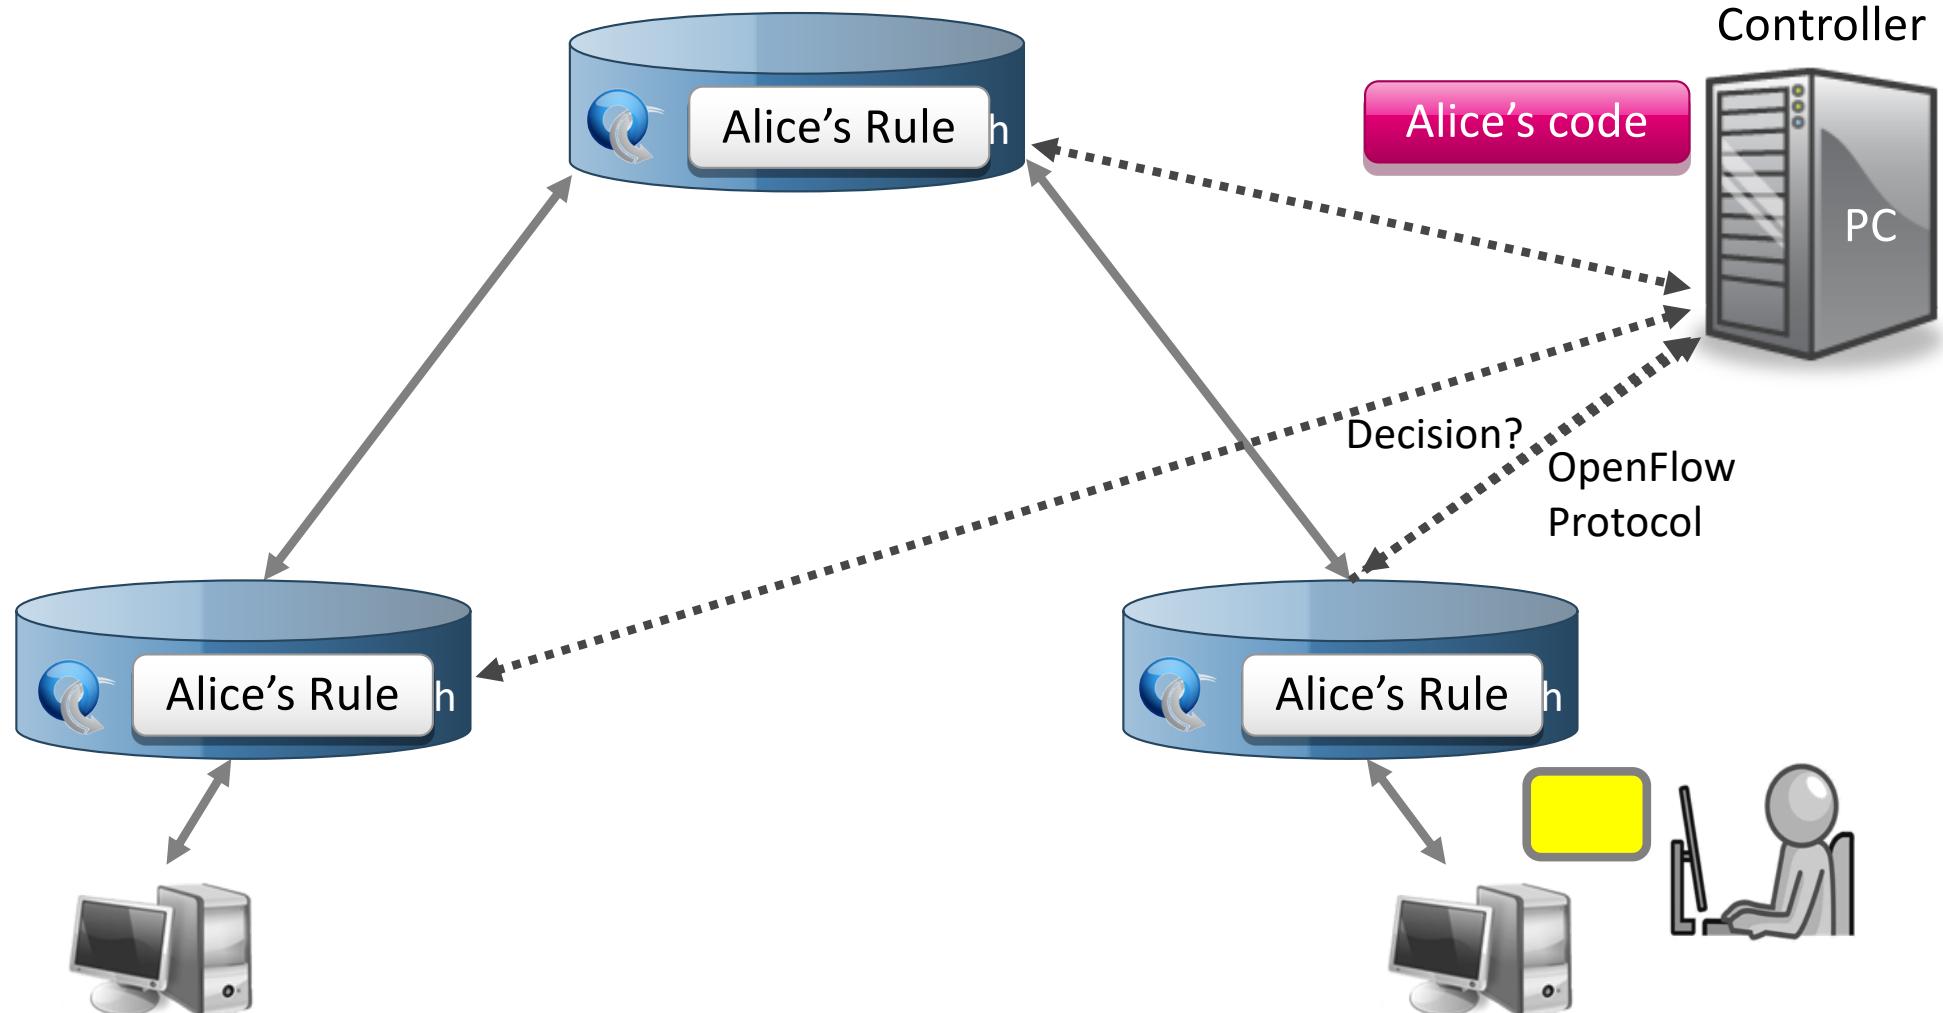
- ▶ Embodiment of policy box
  - Forwarding policy (routing), security policy, energy policy, economic policy, ...
  - Realized in hardware or software
  - Or generic platform running software “policy programs”

# Controller Federation



- ▶ Different boxes belong to different entities
  - Controlled by different controllers
- ▶ Flows route through multiple domains
  - Controllers must coordinate among themselves
  - Possibly exchange limited privileges on datapaths

# OpenFlow Usage



OpenFlow offloads control intelligence to a remote software

# Primitives <Match, Action>

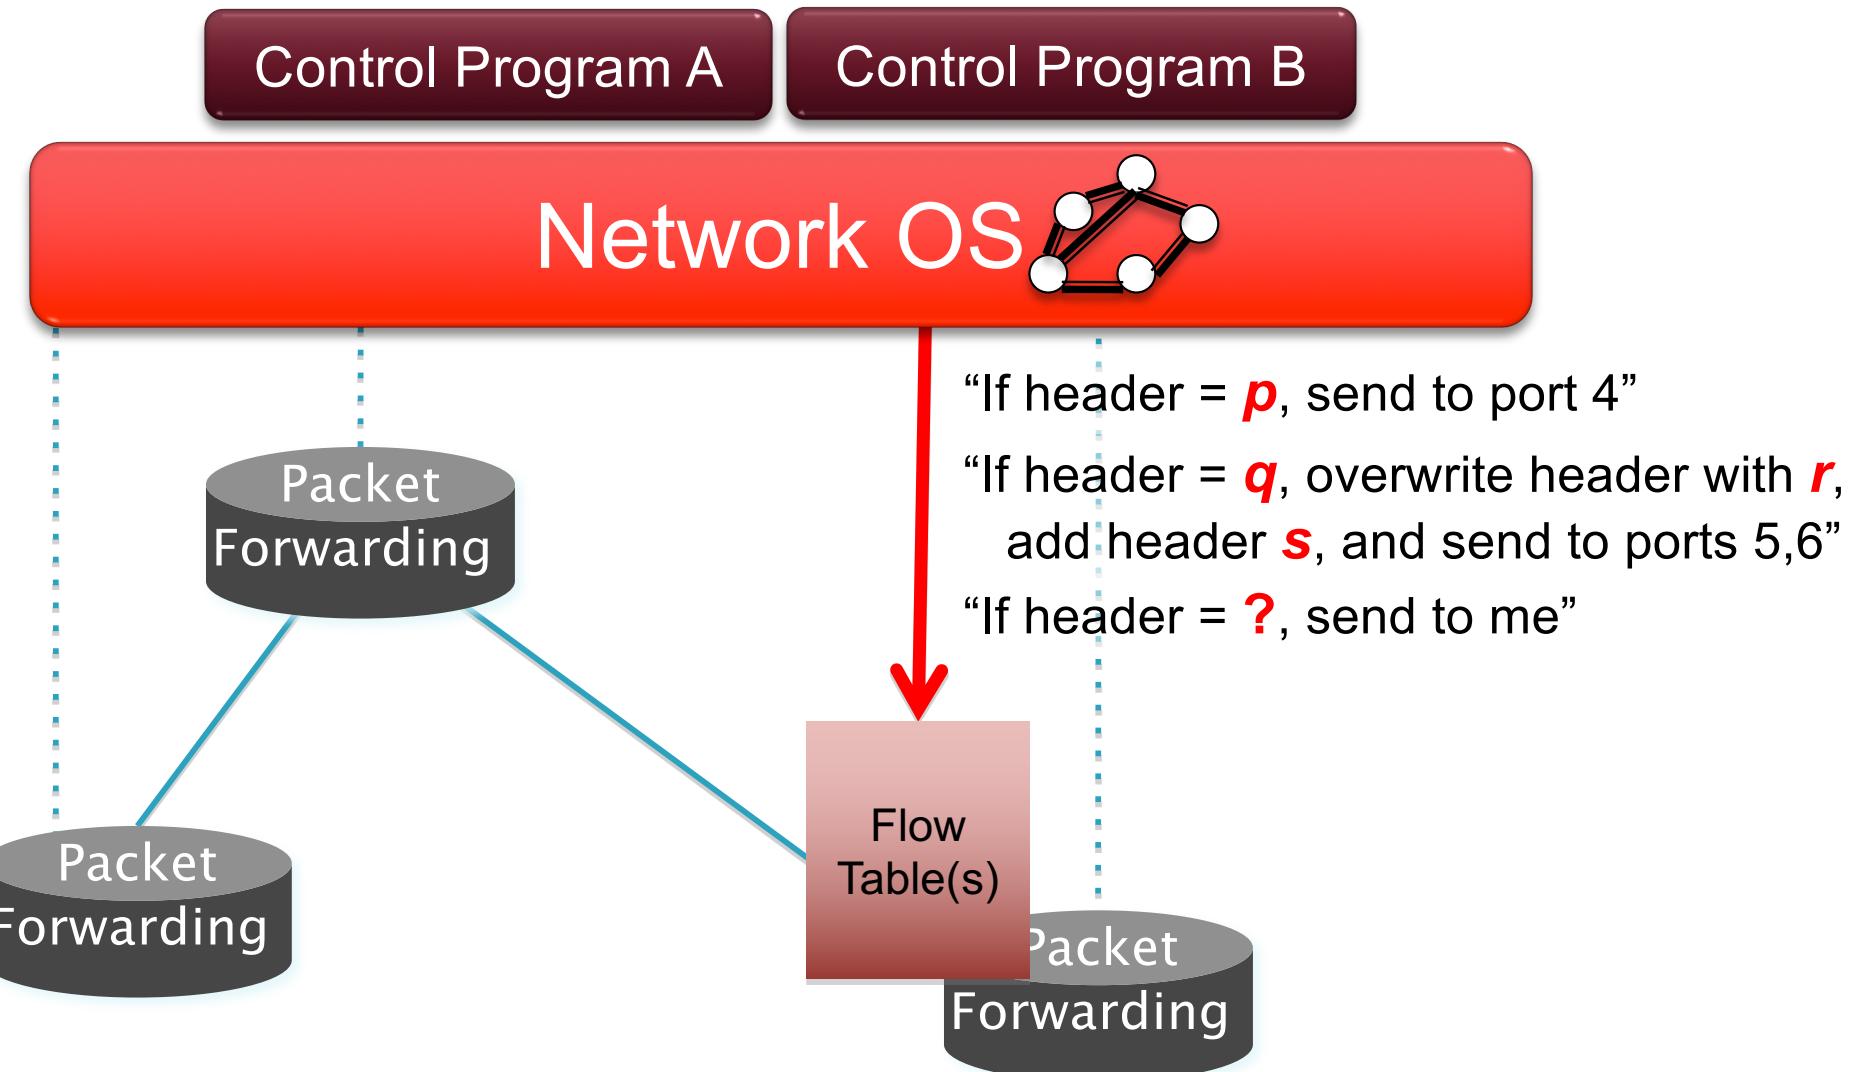
- ▶ **Match** arbitrary bits in headers:



Match: 1000x01xx0101001x

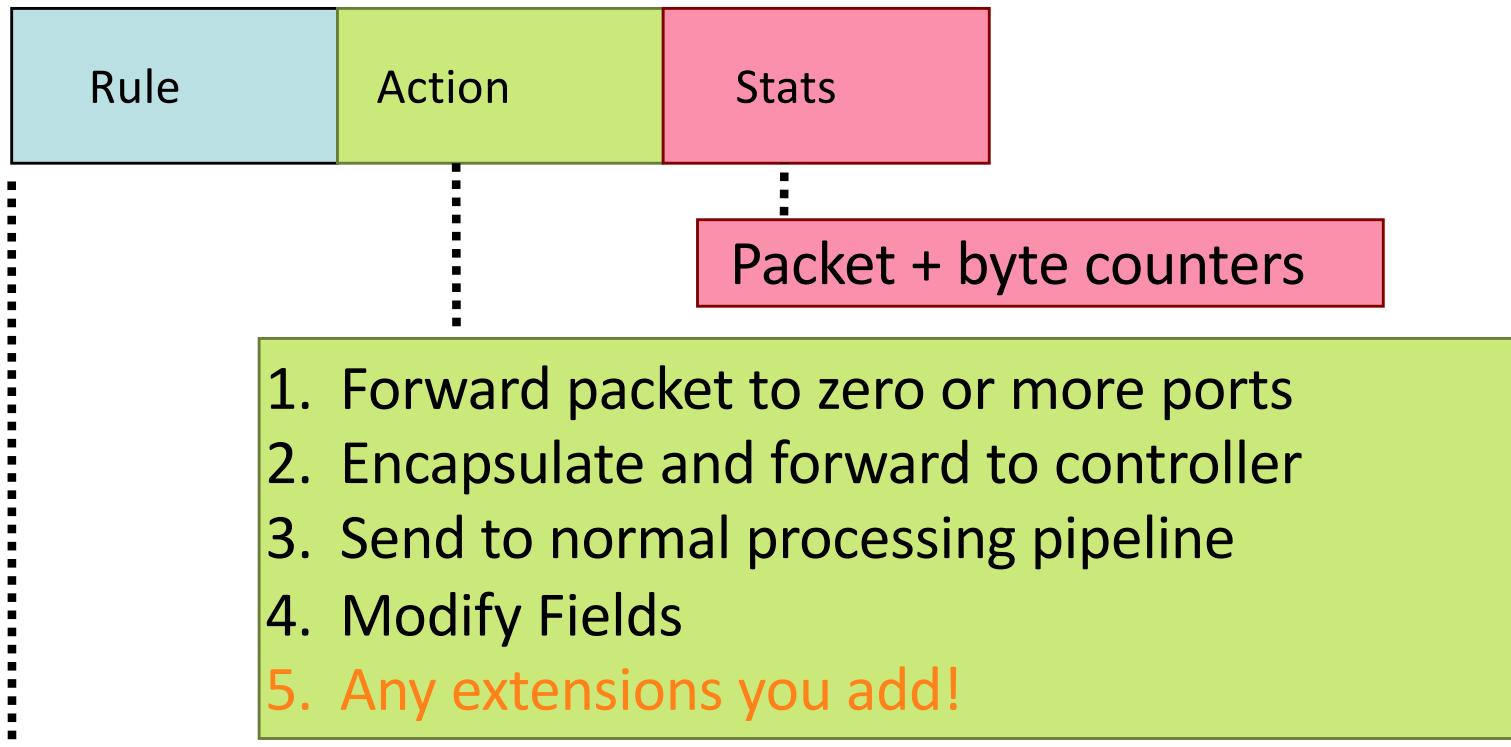
- Match on any header, or new header
  - Allows any flow granularity
- ▶ **Action**
    - Forward to port(s), drop, send to controller
    - Overwrite header with mask, push or pop
    - Forward at specific bit-rate

# OpenFlow Basics



# OpenFlow Basics

## Flow Table Entries



Switch Port	VLAN ID	VLAN pcp	MAC src	MAC dst	Eth type	IP Src	IP Dst	IP ToS	IP Prot	L4 sport	L4 dport
-------------	---------	----------	---------	---------	----------	--------	--------	--------	---------	----------	----------

+ mask what fields to match

# Examples

## Switching

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	*	00:1f...	*	*	*	*	*	*	*	port6

## Flow Switching

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
port3	00:20..	00:1f..	0800	vlan1	1.2.3.4	5.6.7.8	4	17264	80	port6

## Firewall

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	*	*	*	*	*	*	*	*	22	drop

# Examples

## Routing

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	*	*	*	*	*	5.6.7.8	*	*	*	port6

## VLAN Switching

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	*	00:1f..	*	vlan1	*	*	*	*	*	port6, port7, port9

# Example of OpenFlow Instruction Set

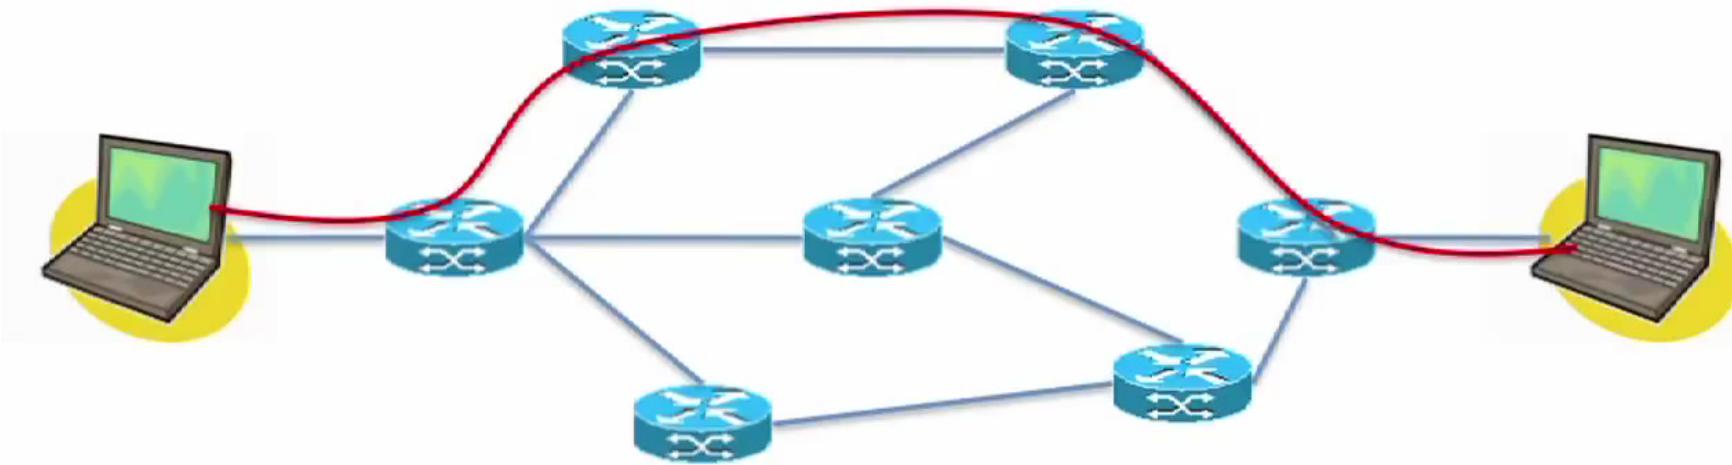
## OpenFlow-enabled Network Device

*Flow Table comparable to an instruction set*

MAC src	MAC dst	IP Src	IP Dst	TCP dport	...	Action	Count
*	10:20:..	*	*	*	*	port 1	250
*	*	*	5.6.7.8	*	*	port 2	300
*	*	*	*	25	*	drop	892
*	*	*	192.*	*	*	local	120
*	*	*	*	*	*	controller	11

# Example: Sophisticated Flow Identification

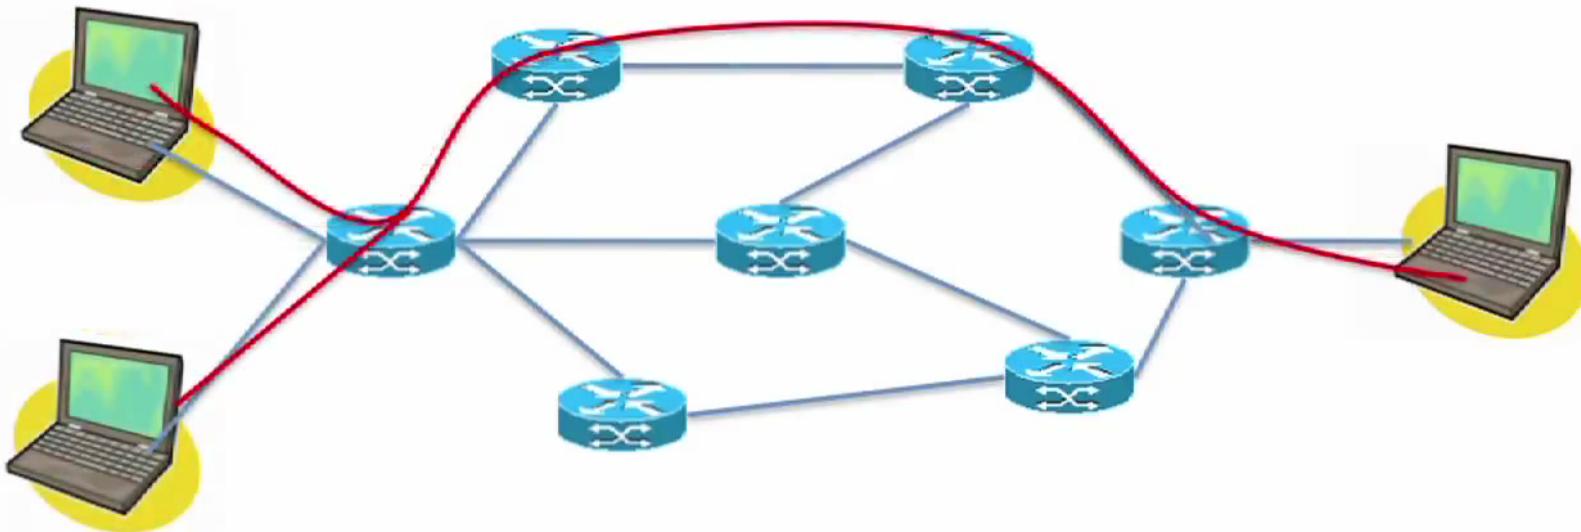
## Application level flow



Switch Port	MAC src	MAC dst	Eth type	VLAN ID	VLAN PCP	IP Src Prefix	IP Dst Prefix	IP Prot	TCP sport	TCP dport
-------------	---------	---------	----------	---------	----------	---------------	---------------	---------	-----------	-----------

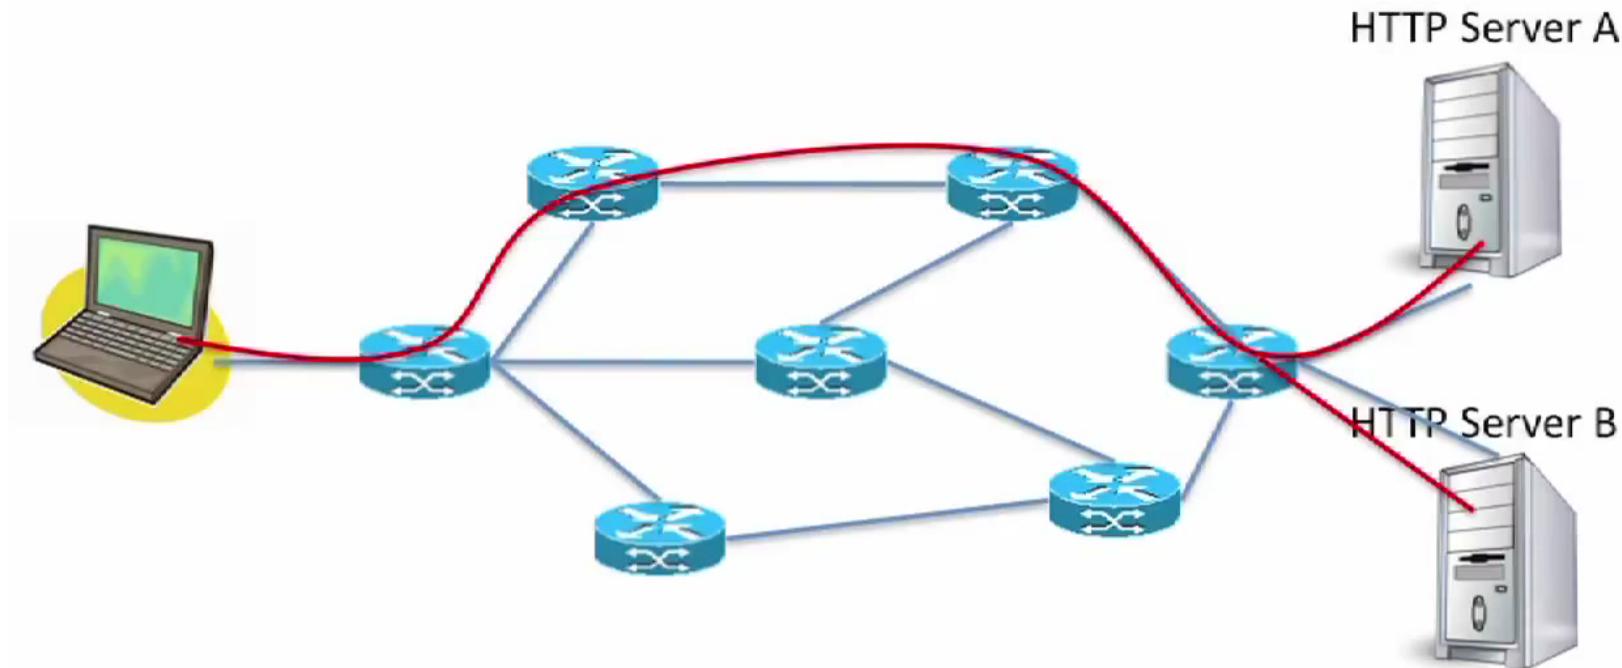
# Example: Sophisticated Flow Identification

IP flow



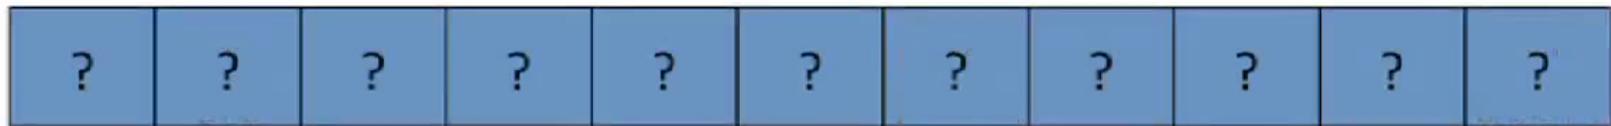
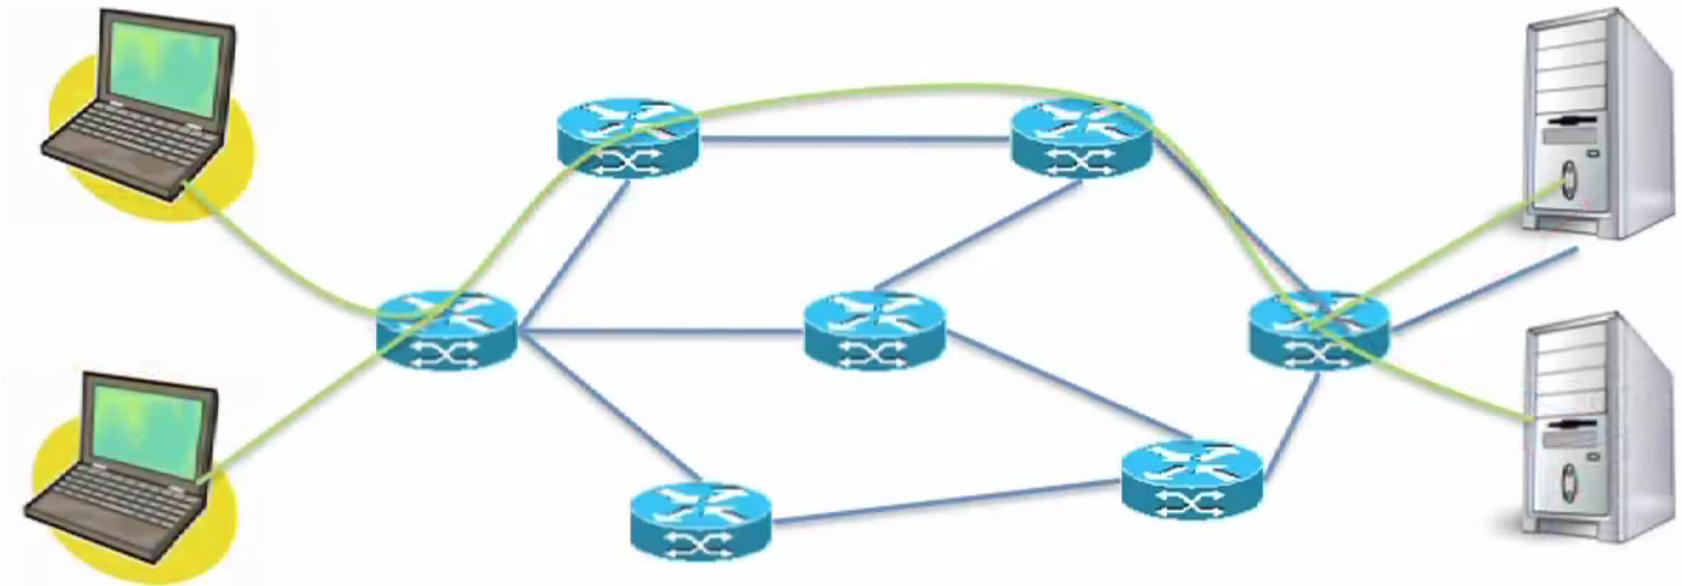
# Example: Sophisticated Flow Identification

## Custom flow



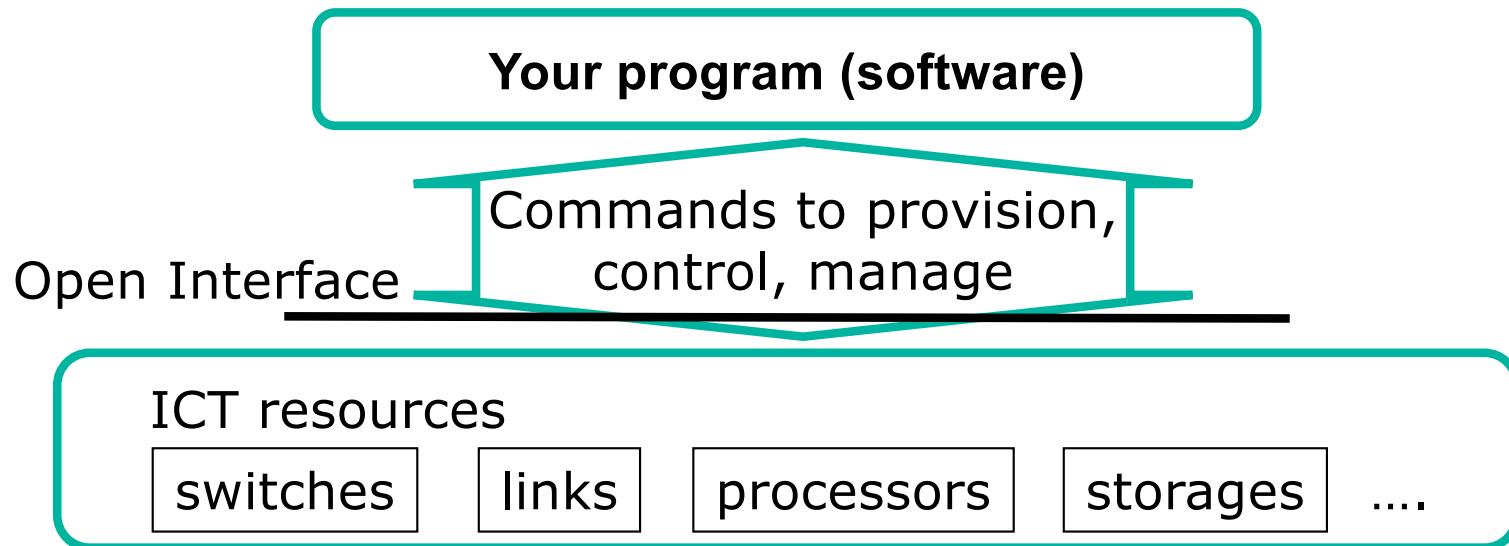
# Example: Sophisticated Flow Identification

My flow



# Software-Defined Networking (SDN)

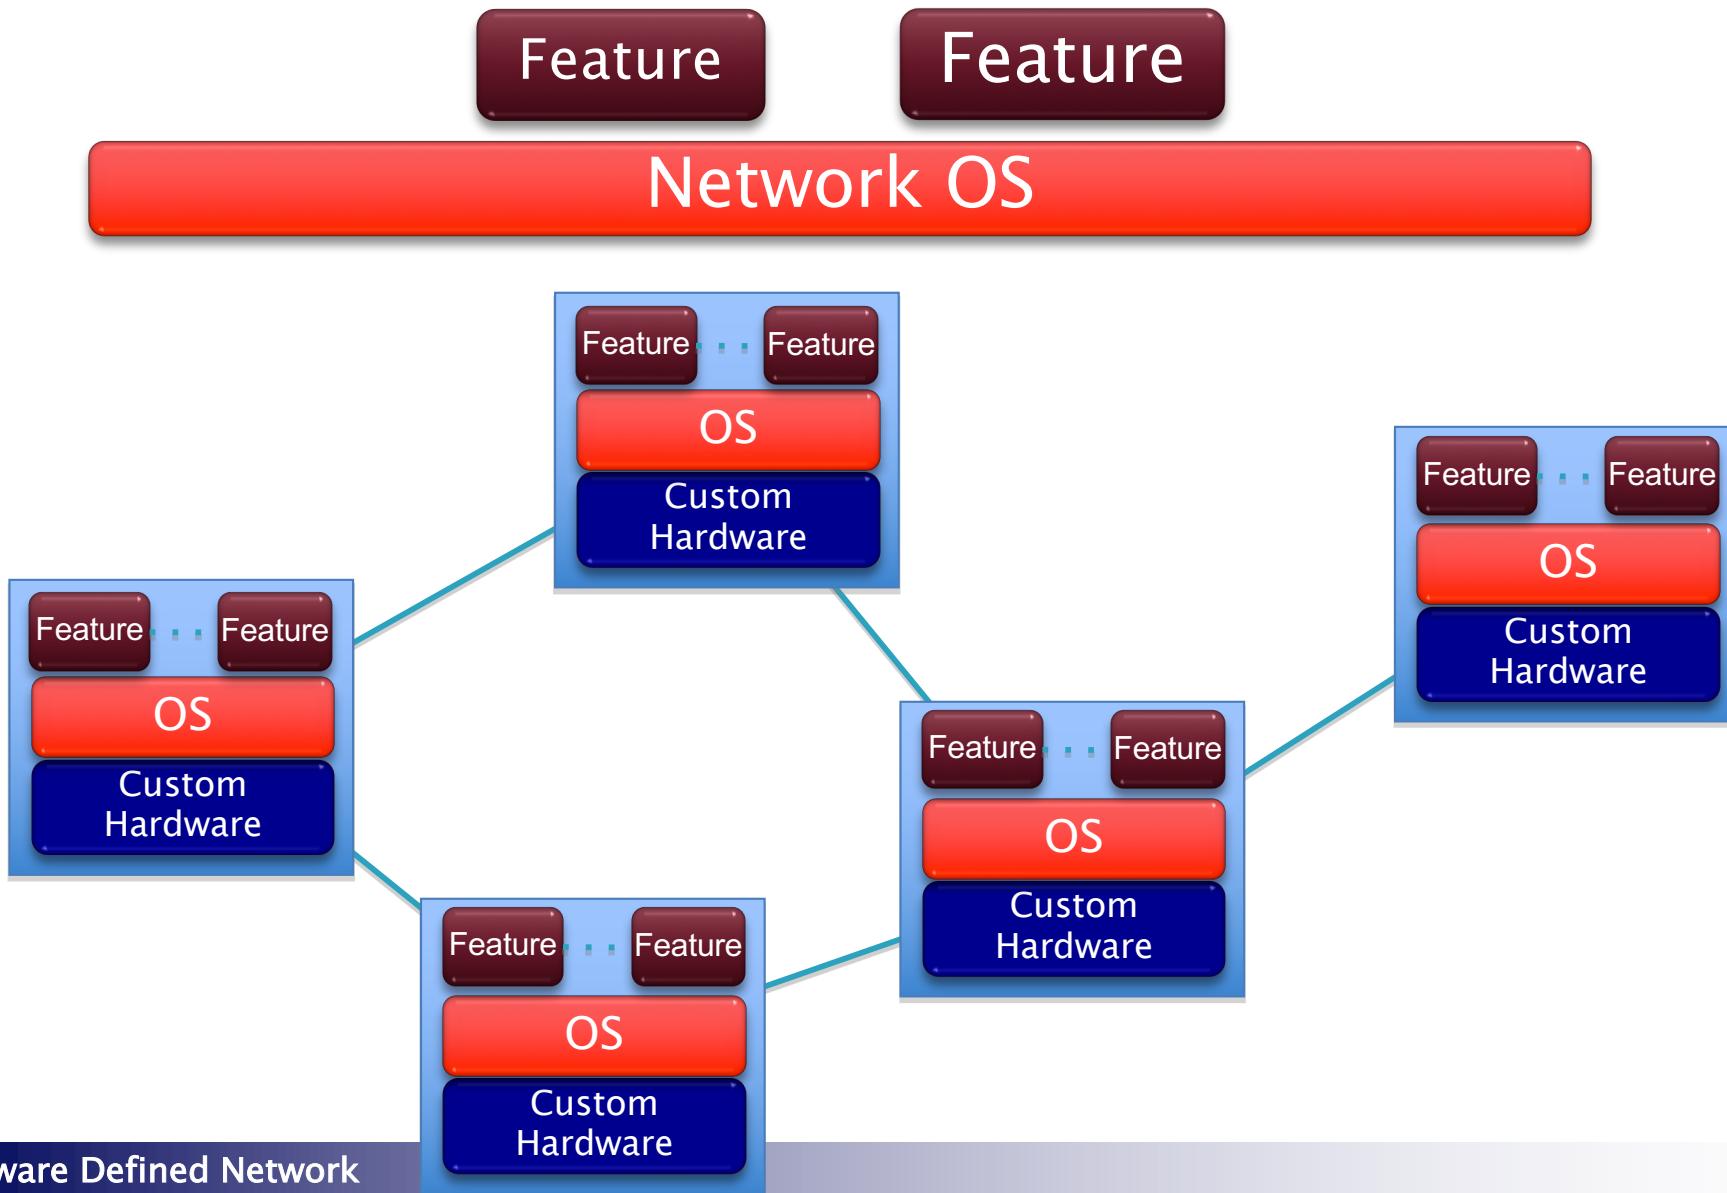
- ▶ Open interface is provided for network resources
- ▶ You define how switches and others behave
  - e.g., if packet header field xxx matches xxx, the packet is forwarded (switch), or discarded (proprietary firewall)
  - e.g., an application notifies your program security breach is found. Your program disable a link (port) to isolate it



# Topics

- ▶ Needs for SDN
- ▶ SDN Principles
- ▶ OpenFlow
- ▶ Impact on Future Networking

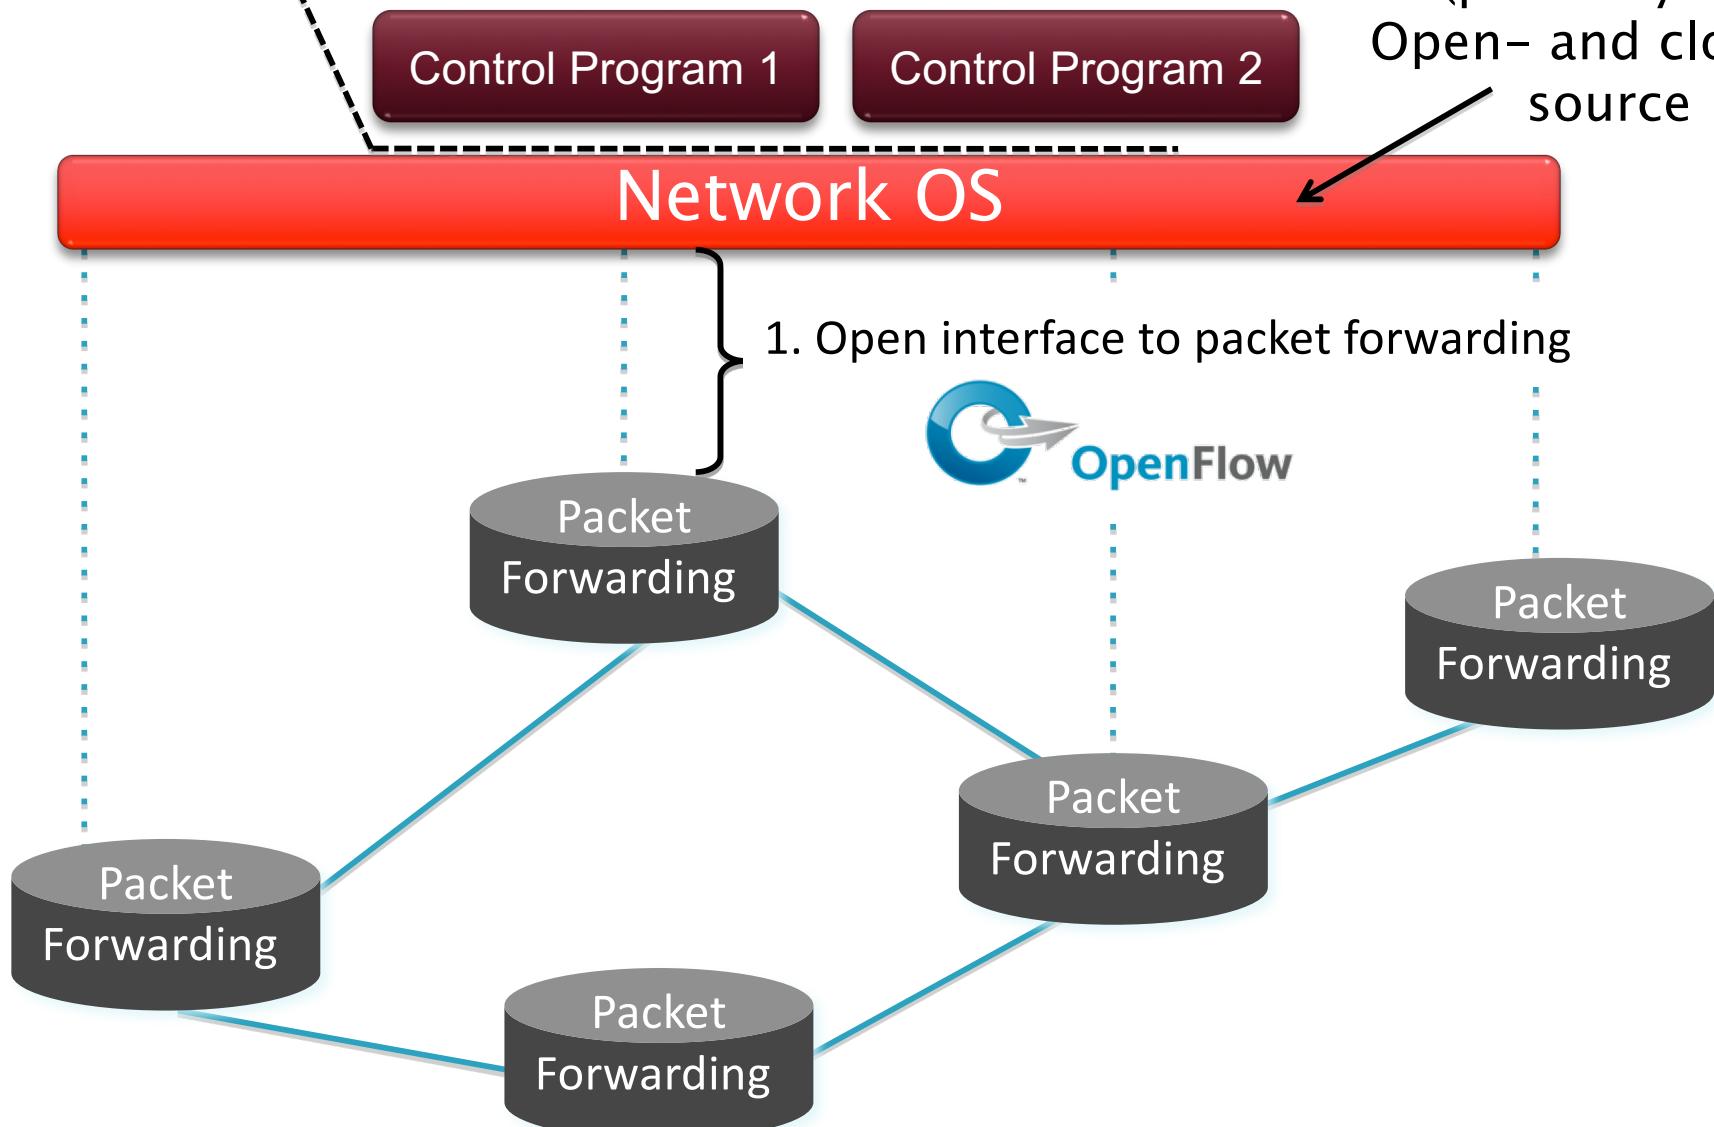
# Recap - The network is changing



# Recap - Software Defined Network (SDN)

3. Consistent, up-to-date global network view

2. At least one Network OS (probably many). Open- and closed-source



# Control/Data Planes Become Separate

- Currently control plane tied to data plane
- NOS runs on servers: observes/controls data plane
- Changes the deployment and business models
  - Can buy the control plane separately from the switches
  - Enabling commodity hardware and 3<sup>rd</sup> party software
- Changes the testing model
  - Simulator to analyze large-scale control planes

# Networking Becomes Edge-Oriented

- Can implement most control functionality at edge
  - Access control, QoS, mobility, migration...
- Network core merely delivers packets edge-to-edge
  - Current protocols do a good job (mostly)
- **Let edge handle all complexity**
  - Complicated matching, actions
  - “Overlay” networking via tunnels
- This has two important implications

# 1. Makes SDN Incrementally Deployable

- Host software often has OpenFlow switch
  - Open vSwitch (OVS) in Linux, Xen,...
- The edge becomes a software switch
  - Core of network can be legacy hardware
- Enables incremental deployment of SDN
  - Might never need OpenFlow in hardware switches....

## 2. Networking Becomes Software-Oriented

- All complicated forwarding done in software (edge)
- And control plane is a program (on a server)...
  - ...not a protocol (on a closed proprietary switch/router)
- We are *programming* the network, not designing it
  - Focus on modularity and abstractions, not packet headers
- **Innovation at software, not hardware, speeds**
- Software lends itself to clean abstractions

# SDN Vision: Networks Become “Normal”

- Hardware: Cheap, interchangeable, Moore’s Law
- Software: Frequent releases, decoupled from HW
- Functionality: Mostly driven by SW
  - Edge (software switch)
  - Control program
- Solid intellectual foundations

# Questions?

- ▶ **Email:** x.che@herts.ac.uk
- ▶ **Office:** LB205A
- ▶ **Tel:** 01707 283206